

TCP/IP Protocol Suite

- A protocol suite for internetworking
- The *TCP/IP Internet Protocols* or, simply, *TCP/IP* is the mostly widely used internetworking protocol suite
- First internetworking protocol suite
- OSI 7-layer model does not explicitly include internetworking
- TCP/IP model replaces the old ISO model

1

TCP/IP Layering

- Layer #5: Application
 - Corresponds to ISO model layers 6 and 7
 - Used for communication among applications
- Layer #4: Transport
 - Corresponds to layer 4 in the ISO model
 - Provides reliable delivery of data
- Layer #3: Internet
 - Defines uniform format of packets forwarded across networks of different technologies
 - Rules for forwarding packets in routers
- Layer #2: Network Interface
 - Corresponds to layer 2 in the ISO model
 - Defines formats for carrying packets in hardware frames
- Layer #1: Hardware
 - Corresponds to layer 1 in the ISO model
 - Defines basic networking hardware

Application	5
Transport	4
Internet	3
Network Interface	2
Physical	1


2

TCP/IP vs OSI

5	Application	Application	7
		Presentation	6
		Session	5
4	Transport	Transport	4
3	Internet	Network	3
2	Network Interface	Data Link	2
1	Physical	Physical	1

3

The Internet



4

Addresses for the Virtual Internet

- One key aspect for creating a virtual network is a single, uniform addressing format
- Can't use hardware addresses because different technologies have different address formats
- Address format must be independent of any particular hardware address format
- Sending host puts destination internet address in packet
- Destination address can be interpreted by any intermediate router
- Routers examine address and forward packet on to the destination
- TCP/IP addresses:
 - Addressing in TCP/IP is specified by the *Internet Protocol* (IP)
 - Each host (not really) is assigned a 32-bit number
 - Called the *IP address* or *Internet address*
 - Unique across entire Internet

5

IP Address Hierarchy

- Each IP address is divided into a prefix and a suffix
- Prefix identifies the network to which a computer is attached
- Suffix identifies a computer within that network
- Address format makes routing efficient
- Every network in a TCP/IP internet is assigned a unique *network number*
- Each host on a specific network is assigned a *host number* or *host address* that is unique *within that network*
- Host's IP address is the combination of the network number (prefix) and host address (suffix)

6

IP Address Assignment

- An IP address does not identify a specific computer
- An IP address specifies an *interface*, or network attachment point, *not* a computer
- A computer with multiple network interconnections (e.g., a router) must be assigned one IP address for each connection
- Global authority assigns unique prefix to network
- Local administrator assigns unique suffix to host

7


IP Address Format Design

- IP designers chose 32 ~~bit~~ addresses
 - Allocate some bits for prefix, some for suffix
 - Large prefix, small suffix: many networks, few hosts per network
 - Small prefix, large suffix: few networks, many hosts per network
- Because of a variety of technologies, need to allow for both large and small networks
- Designers chose a compromise- multiple address formats that allow both large and small prefixes
- Each format is called an address *class*

8

Classful IP Addressing

- First four bits of an address determine the class to which the class belongs
- Specify how the remainder of the address is divided into prefix and suffix



9

Dotted Decimal Notation

- Shorthand for IP address
- Allows humans to avoid binary
- Represents each byte in decimal separated by dots
- Four decimal values per 32-bit address
- Each decimal number:
 - Represents eight bits
 - Is between 0 and 255

32-bit Binary Number	Equivalent Dotted Decimal
10000001 00110100 00000110 00000000	129 . 52 . 6 . 0
11000000 00000101 00110000 00000011	192 . 5 . 48 . 3
00001010 00000010 00000000 00100101	10 . 2 . 0 . 37
10000000 00001010 00000010 00000011	128 . 10 . 2 . 3
10000000 10000000 11111111 00000000	128 . 128 . 255 . 0

10

Classes and Dotted Decimal Notation

- Class identified by the decimal value of the first byte

Class	Range
A	0 - 127
B	128 - 191
C	192 - 223
D	224 - 239
E	240 - 255

11


Classful Addressing and Network Sizes

- Maximum network size determined by class of address:
 - Class A large
 - Class B medium
 - Class C small

Address Class	Bits In Prefix	Maximum Number of Networks	Bits In Suffix	Maximum Number Of Hosts Per Network
A	7	128	24	16777216
B	14	16384	16	65536
C	21	2097152	8	256

12

Classful Addressing Example



13

Address Masks

- IP address space was being exhausted
- Since one of three possible sizes had to be chosen many addresses were unused
- Idea: Instead of having three distinct classes, allow the division between prefix and suffix to occur on an arbitrary bit boundary
- Two pieces of information must be kept for each address:
 - The 32-bit address itself
 - A 32-bit value that specifies the boundary between the prefix and the suffix (*address mask* or *subnet mask*)
 - * 1 bits mark the network prefix
 - * 0 bits mark the host suffix


14

Address Masks Example (1/2)

- Suppose we have the 32-bit destination address 128.10.2.3
- 10000000 00001010 00000010 00000011
- And a 32-bit mask 255.255.0.0
- 11111111 11111111 00000000 00000000
- In order a router to find the correct subnet to route an incoming packet to, performs a *logical and* between the destination address and the mask
- 10000000 00001010 00000000 00000000
- Which is 128.10.0.0 in the dotted decimal notation
- This is compared against a value in the router's routing table

15

Address Masks Example (2/2)



16

Classless Inter-Domain Routing

- A dotted decimal notation for address masks
 - The mask associated with an address is specified by appending a slash to an address and the size of the mask in decimal
- Example:
 - Classful: 128.10.0.0 (16-bit prefix, 16-bit suffix)
 - CIDR: 128.10.0.0/16

17

Address Translation

- Upper levels of protocol stack use *protocol addresses*
- Network hardware must use *hardware address* for eventual delivery
- Protocol address must be translated into hardware address for delivery; will discuss three methods
- Upper levels use only protocol addresses
 - ``Virtual network'' addressing scheme
 - Hides hardware details
- Translation occurs at *data link layer*
 - Upper layer hands down protocol address of destination
 - Data link layer translates into hardware address for use by hardware layer

18

Address Resolution

- Finding hardware address for protocol address:
 - *Address resolution*
 - Data link layer *resolves* protocol address to hardware address
- Resolution is local to a network
- Network component only resolves address for other components on same network

19

Address Resolution Example

- A resolves protocol address for B for protocol messages from an application on A sent to an application on B
- A does *not* resolve a protocol address for F
- Through the internet layer, A delivers to F by routing through R1 and R2
- A resolves R1 hardware address
- Network layer on A passes packet containing destination protocol address F for delivery to R1



20

Address Resolution Techniques

- Association between a protocol address and a hardware address is called a *binding*
- Three techniques:
 - Table lookup:
 - * Bindings stored in memory with protocol address as key
 - * Data link layer looks up protocol address to find hardware address
 - Closed-form computation:
 - * Protocol address based on hardware address
 - * Data link layer derives hardware address from protocol address
 - Dynamic:
 - * Network messages used for "just-in-time" resolution
 - * Data link layer sends message requesting hardware address; destination responds with its hardware address

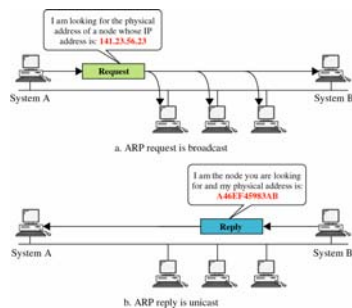
21

Address Resolution Protocol (ARP)

- IP uses dynamic (plus table lookup as an optimization)
- *Address Resolution Protocol* (ARP) - part of TCP/IP protocol suite
- Two-part protocol
 - Request from source asking for hardware address
 - Reply from destination carrying hardware address
- ARP request message dropped into hardware frame and broadcast
- Uses separate protocol type in hardware frame (Ethernet: 0x0806)
- Sender inserts IP address into message and broadcast
- Every other computer examines request
- Computer whose IP address is in request responds
- Puts hardware address in response
- Unicasts to sender
- Original requester can then extract hardware address and send IP packet to destination

22

ARP Operation



23

ARP Packet

- The target machine responds by filling in the missing address
- Swaps the target and sender pairs and changes the operation to a reply
- A will add an entry into its *ARP cache*
- Other nodes may note the IP to MAC address mapping of the node that sent the ARP request

Hardware Type		Protocol Type
Hardware length	Protocol length	Operation
		Request 1, Reply 2
Sender hardware address (For example, 6 bytes for Ethernet)		
Sender protocol address (For example, 4 bytes for IP)		
Target hardware address (For example, 6 bytes for Ethernet) (It is not filled in a request)		
Target protocol address (For example, 4 bytes for IP)		


24

ARP Cache

- Using ARP for each IP packet adds two packets of overhead for each IP packet
- Computer caches ARP responses
 - Flushes cache at system startup
 - Entries discarded periodically
- Cache searched prior to sending ARP request
- ARP lookup algorithm:
 - Look for target IP address, B, in ARP table
 - If not found
 - * Send ARP request
 - * Receive reply with B's hardware address
 - * Add entry to table
- Return hardware address from table

25

ARP Use Cases



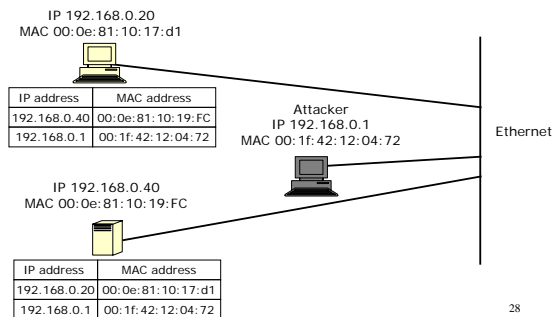
26

ARP Vulnerability

- ARP replies overwrite existing information
- Also, legitimate hosts send ARP replies on joining network or changing IP address
- Not in response to any ARP request
- ARP replies have no proof of origin, so an attacker can claim any hardware address
- Basis for many more attacks
- Known as a *man-in-the-middle* (MITM) attack

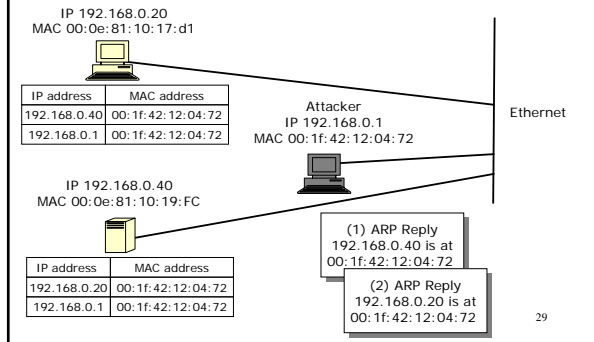
27

Example ARP Attack (1/3)




28

Example ARP Attack (2/3)



Example ARP Attack (3/3)



ARP Attack Effect

- Attacker keeps a *relay index*: a table containing the true association between MAC addresses and IP addresses
- But the two devices at 192.168.0.20 and 192.18.0.40 update their ARP caches with false information
- All traffic for 192.168.0.20 and 192.168.0.40 gets sent to attacker
- Attacker can re-route this traffic to the correct devices using his relay index
- So these devices are oblivious to the attack
- Many implementations out there for many purposes:
 - Capturing traffic in switched Ethernet environments (i.e. active sniffing)
 - Hijacking any application protocol that relies on IP
 - Bypassing firewalls
- * <http://www.althes.fr/ressources/avis/smartspoofing.htm>
- MITM attacks on cryptographically secure protocols
- Denial of service

31