

# TCP/IP Protocol Suite

- A protocol suite for internetworking
- The *TCP/IP Internet Protocols* or, simply, *TCP/IP* is the mostly widely used internetworking protocol suite
- First internetworking protocol suite
- OSI 7-layer model does not explicitly include internetworking
- TCP/IP model replaces the old ISO model

# TCP/IP Layering

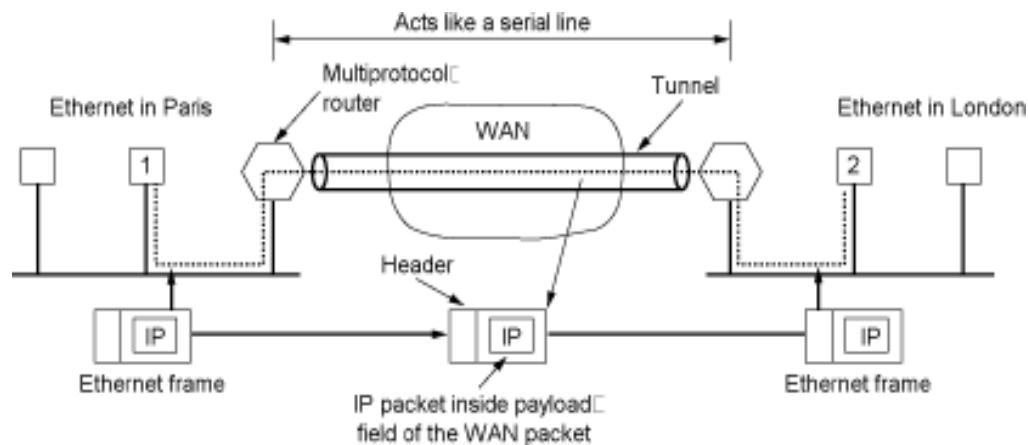

- Layer #5: Application
  - Corresponds to ISO model layers 6 and 7
  - Used for communication among applications
- Layer #4: Transport
  - Corresponds to layer 4 in the ISO model
  - Provides reliable delivery of data
- Layer #3: Internet
  - Defines uniform format of packets forwarded across networks of different technologies
  - Rules for forwarding packets in routers
- Layer #2: Network Interface
  - Corresponds to layer 2 in the ISO model
  - Defines formats for carrying packets in hardware frames
- Layer #1: Hardware
  - Corresponds to layer 1 in the ISO model
  - Defines basic networking hardware

|                   |   |
|-------------------|---|
| Application       | 5 |
| Transport         | 4 |
| Internet          | 3 |
| Network Interface | 2 |
| Physical          | 1 |

# TCP/IP vs OSI

|   |                   |       |              |   |
|---|-------------------|-------|--------------|---|
| 5 | Application       | ..... | Application  | 7 |
|   |                   |       | Presentation | 6 |
|   |                   |       | Session      | 5 |
| 4 | Transport         | ..... | Transport    | 4 |
| 3 | Internet          | ..... | Network      | 3 |
| 2 | Network Interface | ..... | Data Link    | 2 |
| 1 | Physical          | ..... | Physical     | 1 |

# The Internet



# Addresses for the Virtual Internet

- One key aspect for creating a virtual network is a single, uniform addressing format
- Can't use hardware addresses because different technologies have different address formats
- Address format must be independent of any particular hardware address format
- Sending host puts destination internet address in packet
- Destination address can be interpreted by any intermediate router
- Routers examine address and forward packet on to the destination
- TCP/IP addresses:
  - Addressing in TCP/IP is specified by the *Internet Protocol* (IP)
  - Each host (not really) is assigned a 32-bit number
  - Called the *IP address* or *Internet address*
  - Unique across entire Internet

# IP Address Hierarchy

- Each IP address is divided into a prefix and a suffix
- Prefix identifies the network to which a computer is attached
- Suffix identifies a computer within that network
- Address format makes routing efficient
- Every network in a TCP/IP internet is assigned a unique *network number*
- Each host on a specific network is assigned a *host number* or *host address* that is unique *within that network*
- Host's IP address is the combination of the network number (prefix) and host address (suffix)

# IP Address Assignment

- An IP address does not identify a specific computer
- An IP address specifies an *interface*, or network attachment point, *not* a computer
- A computer with multiple network interconnections (e.g., a router) must be assigned one IP address for each connection
- Global authority assigns unique prefix to network
- Local administrator assigns unique suffix to host

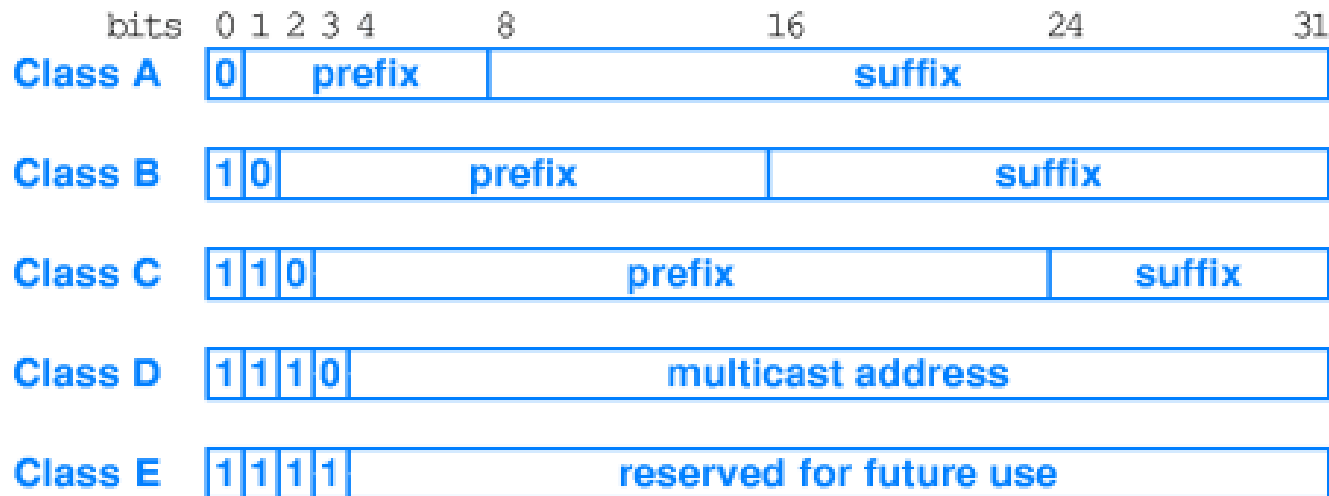
# IP Address Format Design

- IP designers chose 32-bit addresses
  - Allocate some bits for prefix, some for suffix
  - Large prefix, small suffix: many networks, few hosts per network
  - Small prefix, large suffix: few networks, many hosts per network
- Because of a variety of technologies, need to allow for both large and small networks
- Designers chose a compromise - multiple address formats that allow both large and small prefixes
- Each format is called an address *class*



# Classful IP Addressing

- First four bits of an address determine the class to which the class belongs
- Specify how the remainder of the address is divided into prefix and suffix



# Dotted Decimal Notation

- Shorthand for IP address
- Allows humans to avoid binary
- Represents each byte in decimal separated by dots
- Four decimal values per 32-bit address
- Each decimal number:
  - Represents eight bits
  - Is between 0 and 255

| 32-bit Binary Number                | Equivalent Dotted Decimal |
|-------------------------------------|---------------------------|
| 10000001 00110100 00000110 00000000 | 129.52.6.0                |
| 11000000 00000101 00110000 00000011 | 192.5.48.3                |
| 00001010 00000010 00000000 00100101 | 10.2.0.37                 |
| 10000000 00001010 00000010 00000011 | 128.10.2.3                |
| 10000000 10000000 11111111 00000000 | 128.128.255.0             |

# Classes and Dotted Decimal Notation

- Class identified by the decimal value of the first byte


| Class | Range   |
|-------|---------|
| A     | 0-127   |
| B     | 128-191 |
| C     | 192-223 |
| D     | 224-239 |
| E     | 240-255 |

# Classful Addressing and Network Sizes

- Maximum network size determined by class of address:
  - Class A large
  - Class B medium
  - Class C small

| Address Class | Bits In Prefix | Maximum Number of Networks | Bits In Suffix | Maximum Number Of Hosts Per Network |
|---------------|----------------|----------------------------|----------------|-------------------------------------|
| A             | 7              | 128                        | 24             | 16777216                            |
| B             | 14             | 16384                      | 16             | 65536                               |
| C             | 21             | 2097152                    | 8              | 256                                 |

# Classful Addressing Example



# Address Masks

- IP address space was being exhausted
- Since one of three possible sizes had to be chosen many addresses were unused
- Idea: Instead of having three distinct classes, allow the division between prefix and suffix to occur on an arbitrary bit boundary
- Two pieces of information must be kept for each address:
  - The 32-bit address itself
  - A 32-bit value that specifies the boundary between the prefix and the suffix (*address mask* or *subnet mask*)
    - \* 1 bits mark the network prefix
    - \* 0 bits mark the host suffix

# Address Masks Example (1/2)

- Suppose we have the 32-bit destination address 128.10.2.3
  - 10000000 00001010 00000010 00000011
- And a 32-bit mask 255.255.0.0
  - 11111111 11111111 00000000 00000000
- In order a router to find the correct subnet to route an incoming packet to, performs a *logical and* between the destination address and the mask
  - 10000000 00001010 00000000 00000000
  - Which is 128.10.0.0 in the dotted decimal notation
- This is compared against a value in the router's routing table

# Address Masks Example (2/2)

IP Address

|    |   |    |   |    |   |   |
|----|---|----|---|----|---|---|
| 19 | • | 30 | • | 84 | • | 5 |
|----|---|----|---|----|---|---|

Mask

|     |   |     |   |     |   |   |
|-----|---|-----|---|-----|---|---|
| 255 | • | 255 | • | 192 | • | 0 |
|-----|---|-----|---|-----|---|---|

|    |   |    |   |    |   |   |
|----|---|----|---|----|---|---|
| 19 | • | 30 | • | 64 | • | 0 |
|----|---|----|---|----|---|---|

Subnet Address

↓

|     |   |   |   |   |   |   |   |   |
|-----|---|---|---|---|---|---|---|---|
| 84  | 0 | 1 | 0 | 1 | 0 | 1 | 0 | 0 |
| 192 | 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 |
| 64  | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 |



# Classless Inter-Domain Routing

- A dotted decimal notation for address masks
  - The mask associated with an address is specified by appending a slash to an address and the size of the mask in decimal
- Example:
  - Classful: 128.10.0.0 (16-bit prefix, 16-bit suffix)
  - CIDR: 128.10.0.0/16

# Address Translation


- Upper levels of protocol stack use *protocol addresses*
- Network hardware must use *hardware address* for eventual delivery
- Protocol address must be translated into hardware address for delivery; will discuss three methods
- Upper levels use only protocol addresses
  - ``Virtual network'' addressing scheme
  - Hides hardware details
- Translation occurs at *data link* layer
  - Upper layer hands down protocol address of destination
  - Data link layer translates into hardware address for use by hardware layer

# Address Resolution

- Finding hardware address for protocol address:
  - *Address resolution*
  - Data link layer *resolves* protocol address to hardware address
- Resolution is local to a network
- Network component only resolves address for other components on same network

# Address Resolution Example

- $A$  resolves protocol address for  $B$  for protocol messages from an application on  $A$  sent to an application on  $B$
- $A$  does *not* resolve a protocol address for  $F$
- Through the internet layer,  $A$  delivers to  $F$  by routing through  $R1$  and  $R2$
- $A$  resolves  $R1$  hardware address
- Network layer on  $A$  passes packet containing destination protocol address  $F$  for delivery to  $R1$



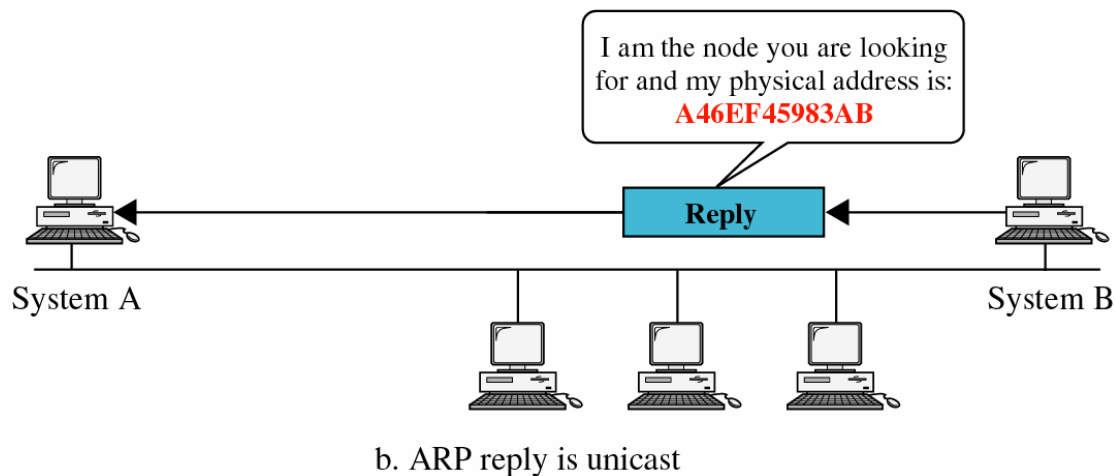
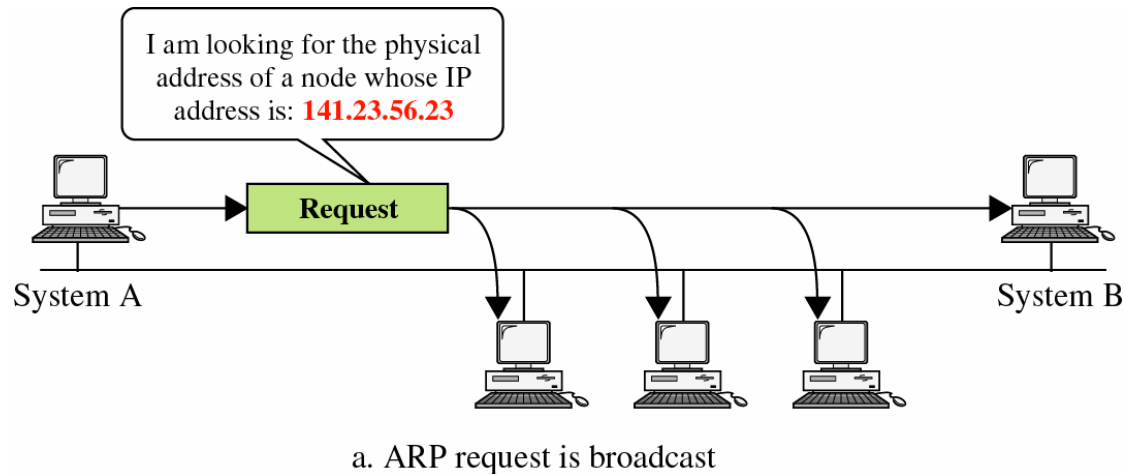
# Address Resolution Techniques

- Association between a protocol address and a hardware address is called a *binding*
- Three techniques:
  - Table lookup:
    - \* Bindings stored in memory with protocol address as key
    - \* Data link layer looks up protocol address to find hardware address
  - Closed-form computation:
    - \* Protocol address based on hardware address
    - \* Data link layer derives hardware address from protocol address
  - Dynamic:
    - \* Network messages used for ``just-in-time'' resolution
    - \* Data link layer sends message requesting hardware address; destination responds with its hardware address

# Address Resolution Protocol (ARP)

- IP uses dynamic (plus table lookup as an optimization)
- *Address Resolution Protocol* (ARP) - part of TCP/IP protocol suite
- Two-part protocol
  - Request from source asking for hardware address
  - Reply from destination carrying hardware address
- ARP request message dropped into hardware frame and broadcast
- Uses separate protocol type in hardware frame (Ethernet: 0x0806)
- Sender inserts IP address into message and broadcast
- Every other computer examines request
- Computer whose IP address is in request responds
  - Puts hardware address in response
  - Unicasts to sender
- Original requester can then extract hardware address and send IP packet to destination

# ARP Operation



# ARP Packet

- The target machine responds by filling in the missing address
  - Swaps the target and sender pairs and changes the operation to a reply
- A will add an entry into its *ARP cache*
  - Other nodes may note the IP to MAC address mapping of the node that sent the ARP request

| Hardware Type   |                 | Protocol Type                   |
|---|-----------------|---------------------------------|
| Hardware length   | Protocol length | Operation<br>Request 1, Reply 2 |
| Sender hardware address<br>(For example, 6 bytes for Ethernet)                                    |                 |                                 |
| Sender protocol address<br>(For example, 4 bytes for IP)  |                 |                                 |
| Target hardware address<br>(For example, 6 bytes for Ethernet)<br>(It is not filled in a request) |                 |                                 |
| Target protocol address<br>(For example, 4 bytes for IP)  |                 |                                 |




# ARP Cache

- Using ARP for each IP packet adds two packets of overhead for each IP packet
- Computer caches ARP responses
  - Flushes cache at system startup
  - Entries discarded periodically
- Cache searched prior to sending ARP request
- ARP lookup algorithm:
  - Look for target IP address, B, in ARP table
  - If not found
    - \* Send ARP request
    - \* Receive reply with B's hardware address
    - \* Add entry to table
  - Return hardware address from table


# ARP Use Cases

Target IP address:  
Destination address in the IP datagram




Case 1. A host has a packet to send to another host on the same network.

Target IP address:  
IP address of a router




Case 2. A host wants to send a packet to another host on another network.  
It must first be delivered to a router.

Target IP address:  
IP address of the appropriate router  
found in the routing table



Case 3. A router receives a packet to be sent to a host on another network.  
It must first be delivered to the appropriate router.

Target IP address:  
Destination address in the IP datagram




Case 4. A router receives a packet to be sent to a host on the same network.

# ARP Vulnerability

- ARP replies overwrite existing information
- Also, legitimate hosts send ARP replies on joining network or changing IP address
  - Not in response to any ARP request
- ARP replies have no proof of origin, so an attacker can claim any hardware address
- Basis for many more attacks
- Known as a *man-in-the-middle* (MITM) attack

# Example ARP Attack (1/3)



# Example ARP Attack (2/3)

IP 192.168.0.20  
MAC 00:0e:81:10:17:d1




| IP address   | MAC address       |
|--------------|-------------------|
| 192.168.0.40 | 00:1f:42:12:04:72 |
| 192.168.0.1  | 00:1f:42:12:04:72 |

Attacker  
IP 192.168.0.1  
MAC 00:1f:42:12:04:72



IP 192.168.0.40  
MAC 00:0e:81:10:19:FC



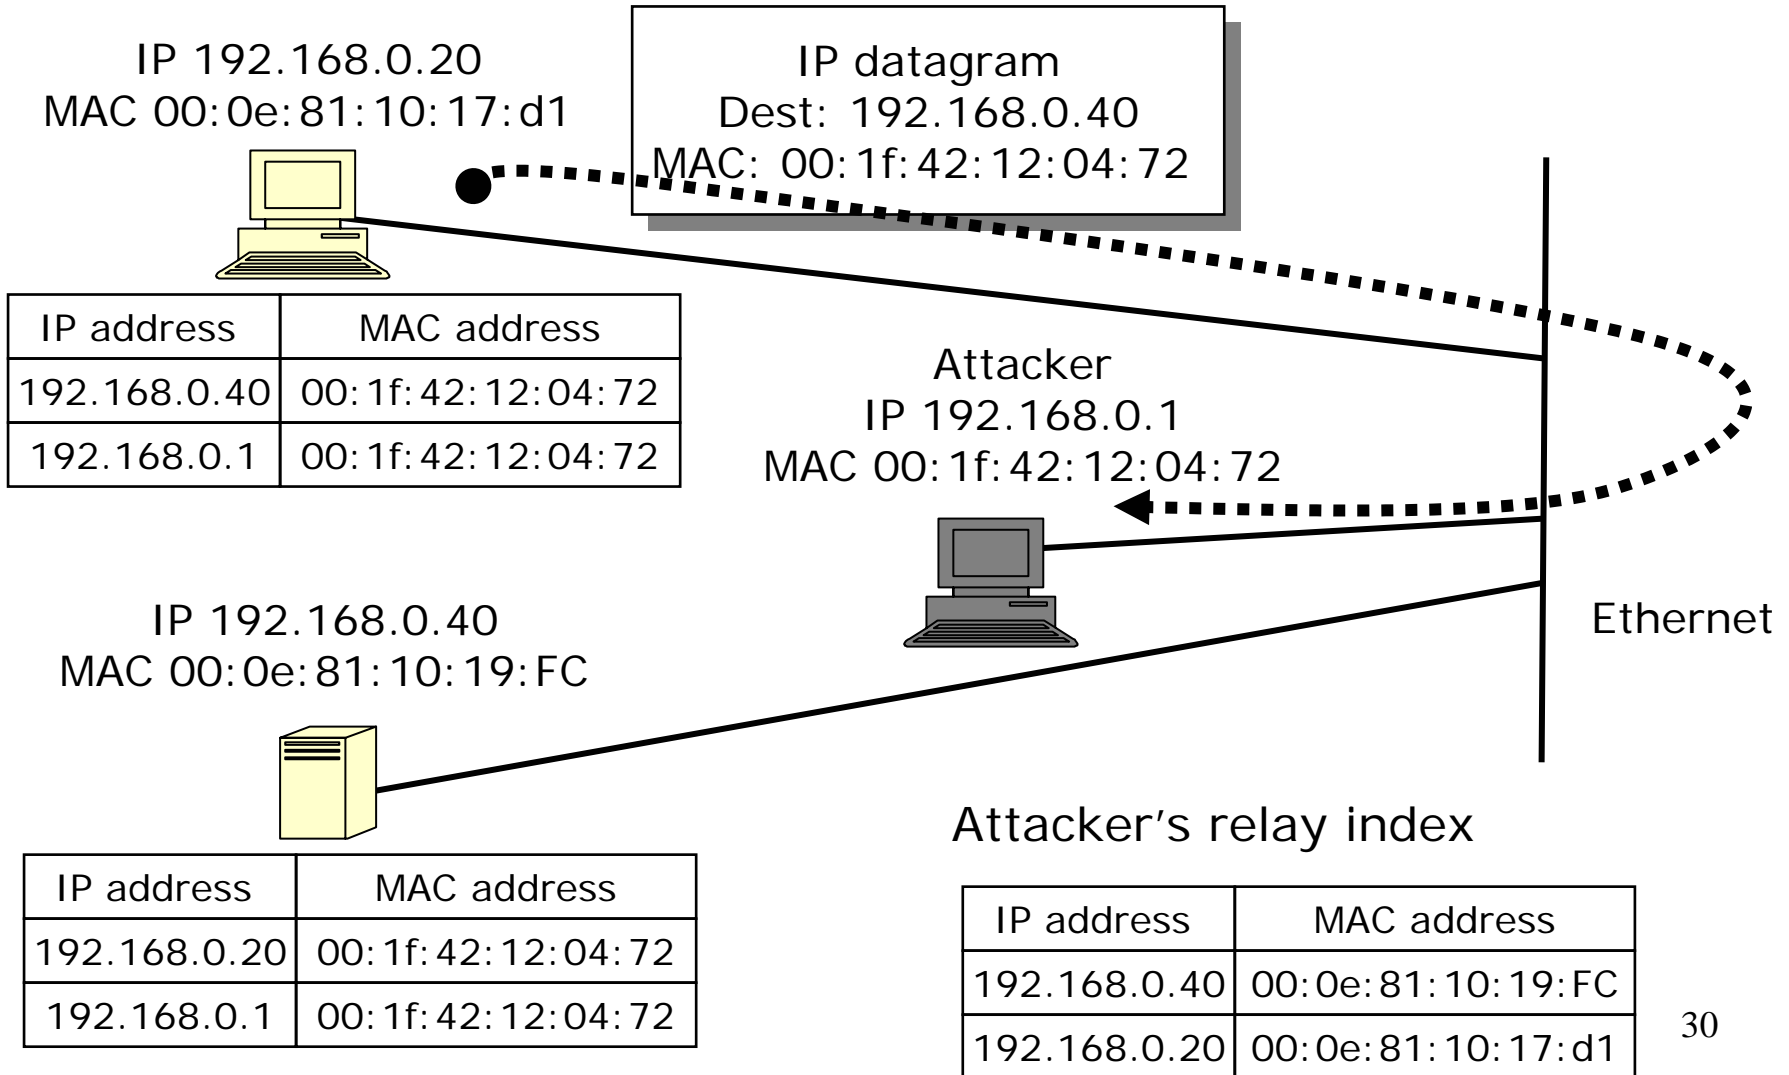
| IP address   | MAC address       |
|--------------|-------------------|
| 192.168.0.20 | 00:1f:42:12:04:72 |
| 192.168.0.1  | 00:1f:42:12:04:72 |

(1) ARP Reply  
192.168.0.40 is at  
00:1f:42:12:04:72

(2) ARP Reply  
192.168.0.20 is at  
00:1f:42:12:04:72

Ethernet

# Example ARP Attack (3/3)



# ARP Attack Effect

- Attacker keeps a *relay index*: a table containing the true association between MAC addresses and IP addresses
- But the two devices at 192.168.0.20 and 192.18.0.40 update their ARP caches with false information
- All traffic for 192.168.0.20 and 192.168.0.40 gets sent to attacker
- Attacker can re-route this traffic to the correct devices using his relay index
- So these devices are oblivious to the attack
- Many implementations out there for many purposes:
  - Capturing traffic in switched Ethernet environments (i.e. active sniffing)
  - Hijacking any application protocol that relies on IP
  - Bypassing firewalls
    - \* <http://www.althes.fr/ressources/avis/smartspoofing.htm>
  - MITM attacks on cryptographically secure protocols
  - Denial of service