

TCP Multiplexing

- TCP socket identified by 4-tuple:
 - Source IP address
 - Source port number
 - Destination IP address
 - Destination port number
- Receiver uses all four values to direct segment to appropriate socket
- A server may support many simultaneous TCP sockets:
 - Each socket identified by its own 4-tuple
- Example:
 - * Web servers have different sockets for each connecting client
 - * Non-persistent HTTP has a different socket for each request

1


Reliable Connection Setup

- Frame, packets and primitives in the network may be delayed for one or another reason and unexpectedly appear in their destination at a time where they are unwanted or unexpected
- This in the connection establishment procedure may cause problems, as a delayed connection request shown up suddenly may result in establishing a connection incorrectly
- To solve this problem, Tomlinson in 1975 introduced the three-way handshake
- This establishment protocol allows each side to begin with a different sequence number
- This way delayed primitives can be directly identified and discarded

2


Three Way Handshake

- Sender executes CONNECT primitive (specifying IP address, port (e.g. 21), the max TCP segment it will accept and other optional user data (e.g. username, password))
- CONNECT primitive sends a TCP segment with the SYN bit on (choosing a sequence number, e.g. 1200) and waits for a response (first arrow)
- When this segment arrives at the receiver, the TCP entity checks to see if there is a process that has done a LISTEN on the port 21 and the gives the incoming segment to that process (i.e. an ftp server)
- If it accepts, an acknowledgement segment is sent back, acknowledging 1200 and announcing its own initial sequence number 4800 (second arrow)
- Finally, sender acknowledges receiver's initial sequence number 4800 and sends data with another initial sequence number
- Then the sequence of TCP segments with data follows



3


Reliable Connection Shutdown



4

Shutdown (Final ACK Lost)


- To cope with lost primitives TCP sets a counter in the receiver when a FIN primitive arrives
- If the final ACK is lost, then the connection is released automatically when it timeouts



5

Shutdown (Request Lost)

- To cope with lost primitives TCP sets a counter in the sender when a FIN primitive is sent
- If the request is lost, then no ACK will ever arrive; a new request is sent automatically when the ACK waiting timeouts



6

TCP Problems and Attacks

- Silly window syndrome (Clark, 1982)
 - If receiver advertises small increases in the receive window then the sender may waste time sending lots of small packets
 - Solution:
 - * Receiver must not advertise small window increases
 - * Increase window by $\min(\text{SMSS}, \text{RecvBuffer}/2)$

7

Nagel's Algorithm

- Small packet problem:
 - Don't want to send a 41 byte packet for each keystroke
 - How long to wait for more data?
- Solution:
 - Allow only one outstanding small (not full sized) segment that has not yet been acknowledged

8

Savage's Attacks

- Congestion control with a misbehaving receiver (greedy Web client)
- Drive a standard TCP sender arbitrarily fast; competing traffic delayed or even discarded
 - ACK division
 - DupACK spoofing
 - Optimistic ACKing
- <http://www.cs.washington.edu/homes/savage/papers/CCR99.pdf>

9


ACK Division (1/2)

- TCP's error control based on byte offsets within a byte stream
- Congestion control implicitly defined in terms of segments rather than bytes
- Attack
 - Upon receiving a segment of N bytes the receiver calculates the resulting ACK
 - Receiver divides this ACK into M, where $M \leq N$, separate ACKs
 - Each covers one of M distinct pieces of the received segment

10

ACK Division (2/2)


- Sender begins with $cwnd=1$
- Incremented for each of the three valid ACKs received
- After one RTT $cwnd=4$ instead of the expected value $cwnd=2$
- Each ACK is valid, covers data that was sent and previously unacknowledged
- Sender's grows $cwnd$ at a rate M faster than usual



11

DupACK Spoofing

- Upon receiving a data segment the receiver sends a long stream of ACKs for the last sequence number received (start of connection: SYN segment)
- $cwnd$ is increased by SMSS for each additional duplicate ACK
- TCP assumes that duplicate ACKs are sent in response to unique and distinct segments




12

Optimistic ACKing

- TCP implicitly assumes that the time between a data segment being sent and an acknowledgement for that segment returning is at least one RTT
- cwnd grown is a function of RTT, sender-receiver pairs with shorter RTTs will transfer data more quickly
- TCP does not use any mechanism to enforce this assumption
- Receiver can emulate a shorter RTT by sending ACKs optimistically for data it has not received yet

13

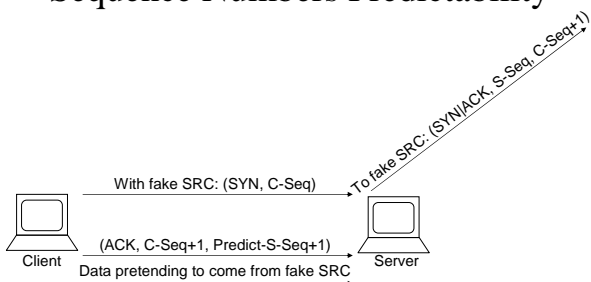
SYN Flooding



- DoS isn't due to traffic volume but to resource exhaustion (memory) in the server
- Early network stacks had a severely limited number of half-open structures available
- Can spoof SRC address with non-existent host
- Solution: SYN cookies – make the SYNACK contents purely a function of SYN contents, therefore, it can be recomputed on reception of next ACK

14

Sequence Numbers Predictability



Particularly dangerous when "fake SRC" is a trusted IP address

15