

# Domain Name System (DNS)

- Internet communication requires IP addresses
- Humans prefer to use computer names
- Automated system available to translate names to addresses
- Known as *Domain Name System (DNS)*
- November 1987
  - RFC 1034: Informational
  - RFC 1035: Implementation details

# DNS Functionality

- Given
  - Name of a computer
- Return
  - Computer's internet (IP) address
- Method
  - Distributed lookup
  - Client contact server(s) as necessary

# Domain Name Syntax

- Alphanumeric segments separated by dots
- Examples:
  - www.tcd.ie
  - ntrg.cs.tcd.ie
  - www.research.att.com
- Most significant part on the right

# Domain Name Acquisition

- Organization
  - Chooses a desired name
  - Must be unique
  - Registers with central authority
  - Placed under one *top-level domain*
- Names subject to international law
  - Trademarks
  - Copyright

# Top-Level Domains (TLDs)

- .com commercial organization
- .edu U.S. educational institution
- .gov U.S. government organization
- .mil U.S. military group
- .net major network provider or other
- .org organization other than above
- .arpa temporary ARPA domain (still used)
- .int international organization
- country code A country (e.g. ie or gr (grrr))

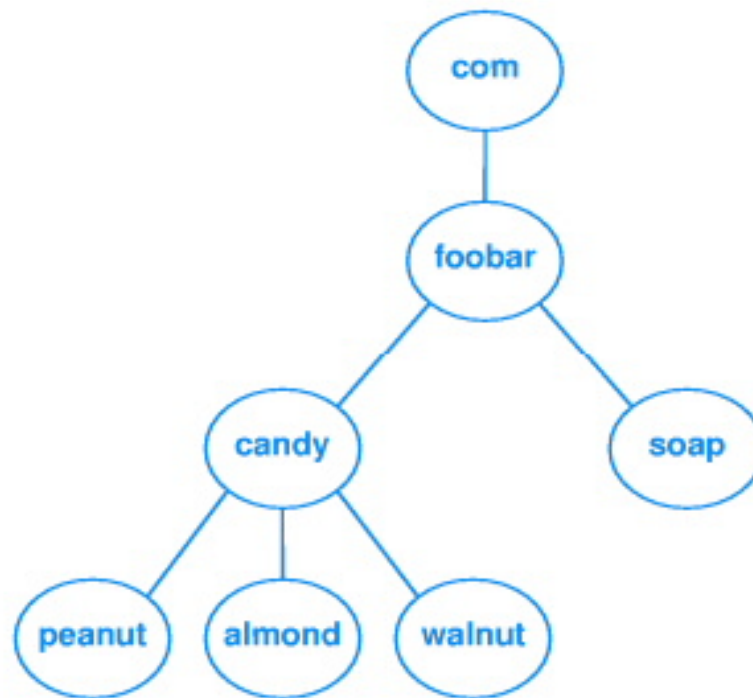
# Within Top-Level Domains

- Subdivision possible
- Arbitrary levels possible
- Not standardized
- Controlled locally by organization

# Example Name Structure

- First level: `com`
- Second level is company name: `cisco`
- Third level is division within company: `security`
- Fourth level either
  - Company subdivision: `crypto`
  - Individual computer: `smtp`

# DNS Illustrated





# DNS Key Concepts

- The number of segments in a domain name corresponds to the naming hierarchy
- There is no universal standard for this hierarchy; each organization can choose its own naming convention
- Furthermore, names within an organization do not need to follow a uniform pattern; individual groups within the organization can choose a hierarchical structure that is appropriate for that group

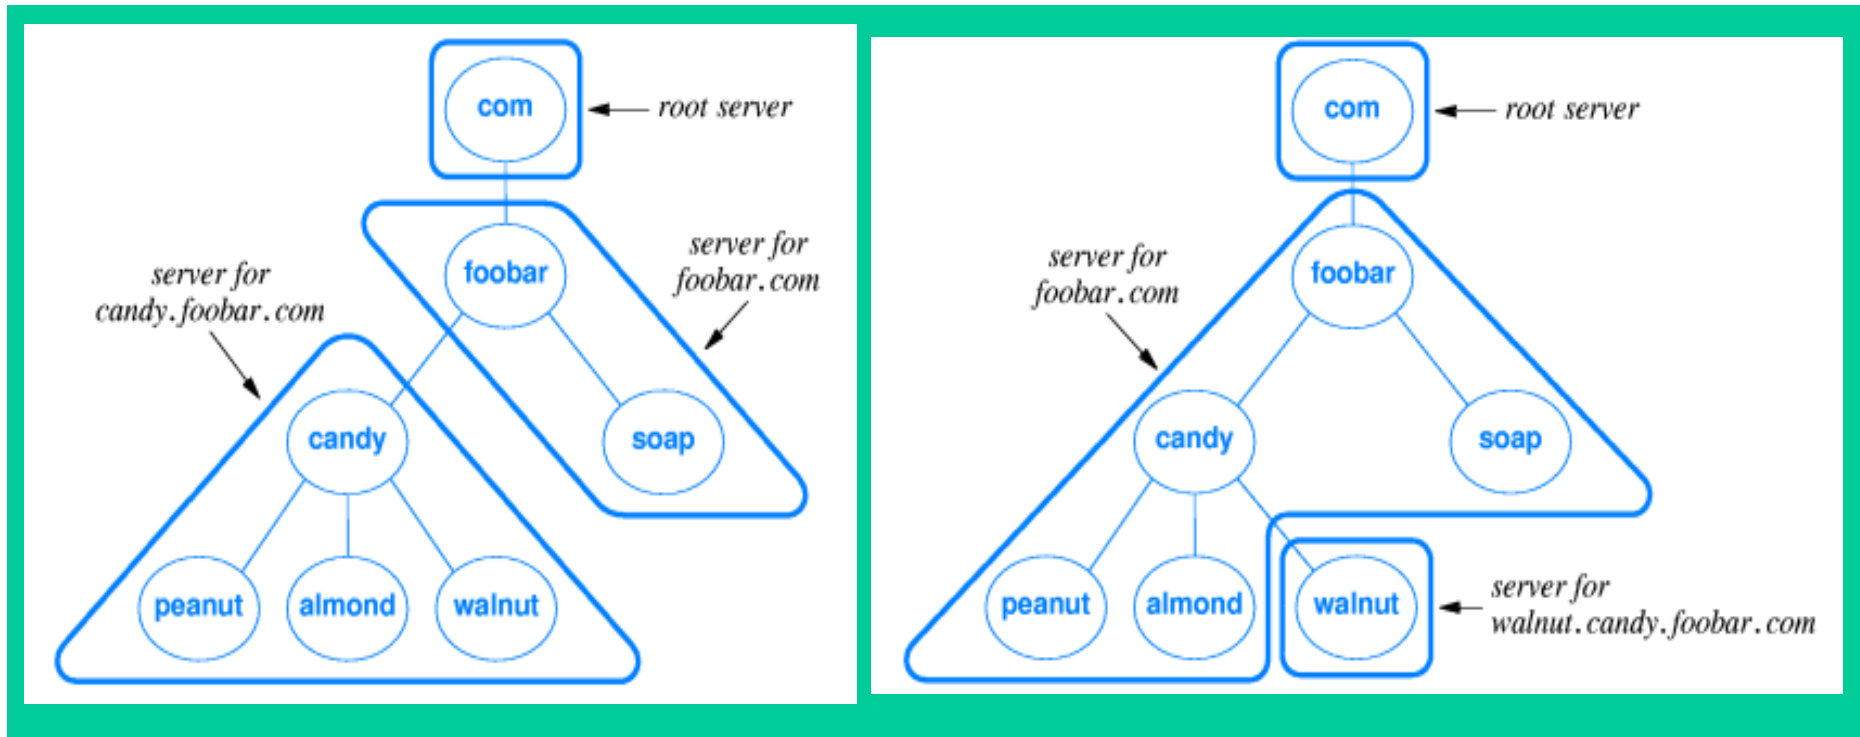
# DNS Client/Server Interaction

- Client known as *resolver*
  - Actually a library that applications link against
- Multiple DNS servers used
- Arranged in a hierarchy
- Each server corresponds to an adjacent part of the global naming hierarchy

# Inter-Server Links


- All domain name servers are linked together to form a unified system
- Each server knows how to reach a root server, and
- How to reach servers that are authorities for names further down the hierarchy

# DNS Hierarchy



- Choice made by organization foobar

# Distributed Hierarchical Database



- Root servers and TLD servers typically do not contain hostname to IP mappings
- They contain mappings for locating authoritative servers

# DNS Root Name Servers

- Contacted by local name server that can not resolve name
- Root name server:
  - Contacts authoritative name server if name mapping not known
  - Gets mapping
  - Returns mapping to local name server

# TLD and Authoritative Servers

- Top-level domain (TLD) servers:  
Responsible for com, org, net, edu, etc., and all top-level country domains ie, gr, ...
- Authoritative DNS servers: Organization's DNS servers, providing authoritative hostname to IP mappings for organization's servers (e.g., web and mail)
  - Can be maintained by organization or service provider

# Local Name Server

- Each ISP (residential ISP, company, university) has one
  - Also called ``default name server''
- When a host makes a DNS query, query is sent to its local DNS server
  - Acts as a proxy, forwards query into hierarchy
  - Reduces lookup latency for commonly searched hostnames



# Caching and Updating Records

- Once (any) name server learns mapping, it *caches* mapping
  - Cache entries timeout (disappear) after some time
  - TLD servers typically cached in local name servers
    - \* Thus root name servers not often visited
- Update/notify mechanisms under design by IETF
  - RFC 2136
  - <http://www.ietf.org/html.charters/dnsind-charter.html>

# DNS Records

DNS: Distributed DB storing resource records (RR)

RR format: (name, value, type, ttl)

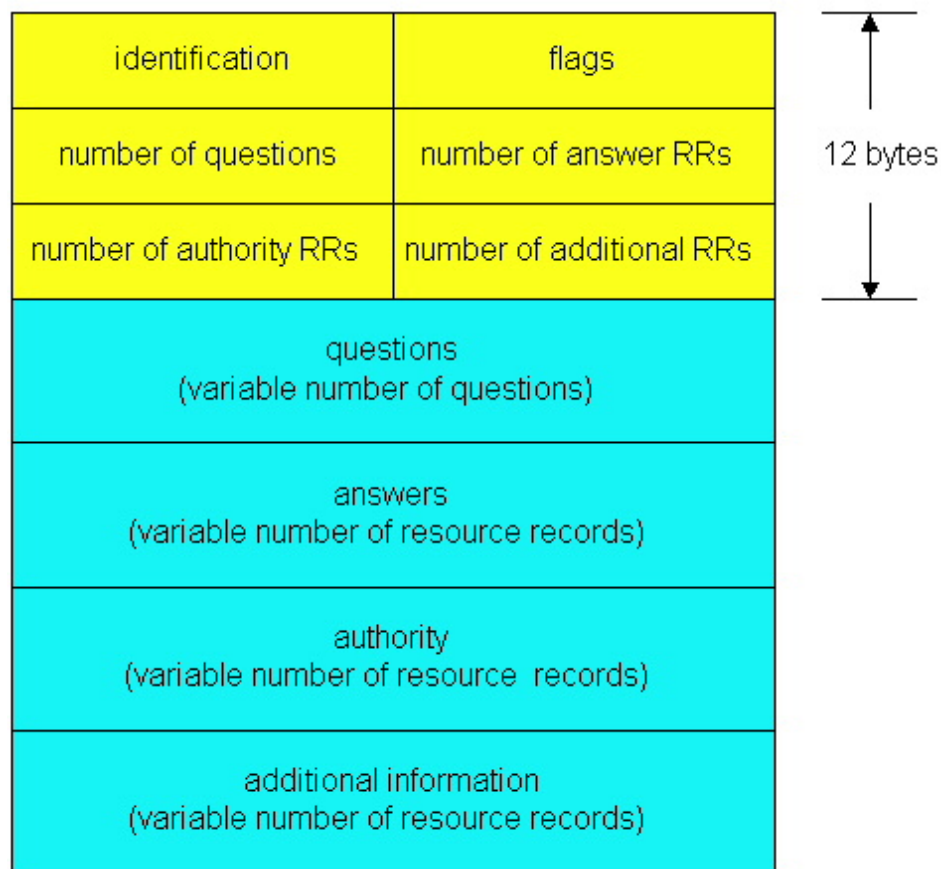
- Type=A
  - **name** is hostname
  - **value** is IP address
- Type=NS
  - **name** is domain (e.g. foo.com)
  - **value** is IP address of authoritative name server for this domain
- Type=CNAME
  - **name** is alias name for some ``canonical'' (the real) name  
www.ibm.com is really  
servereast.backup2.ibm.com
  - **value** is canonical name
- Type=MX
  - **value** is name of mail server associated with **name**

# DNS Protocol and Messages (1/2)

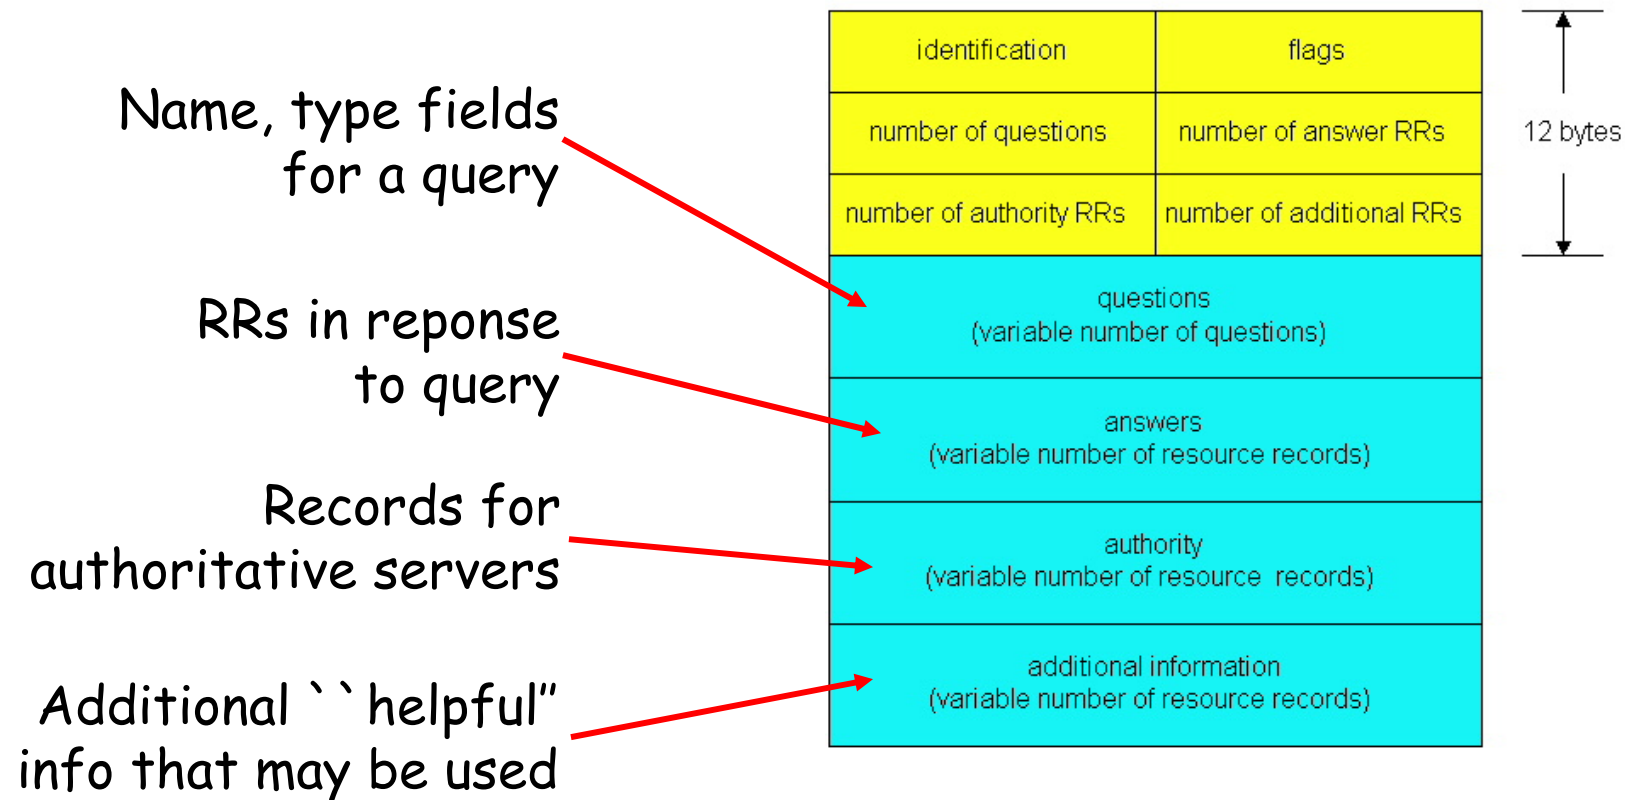
DNS protocol: *Query* and *reply* messages, both with same *message format*

## Message header:

- Identification: 16 bit # for query, reply to query uses same #
- Flags:
  - Query or reply
  - Reply is authoritative



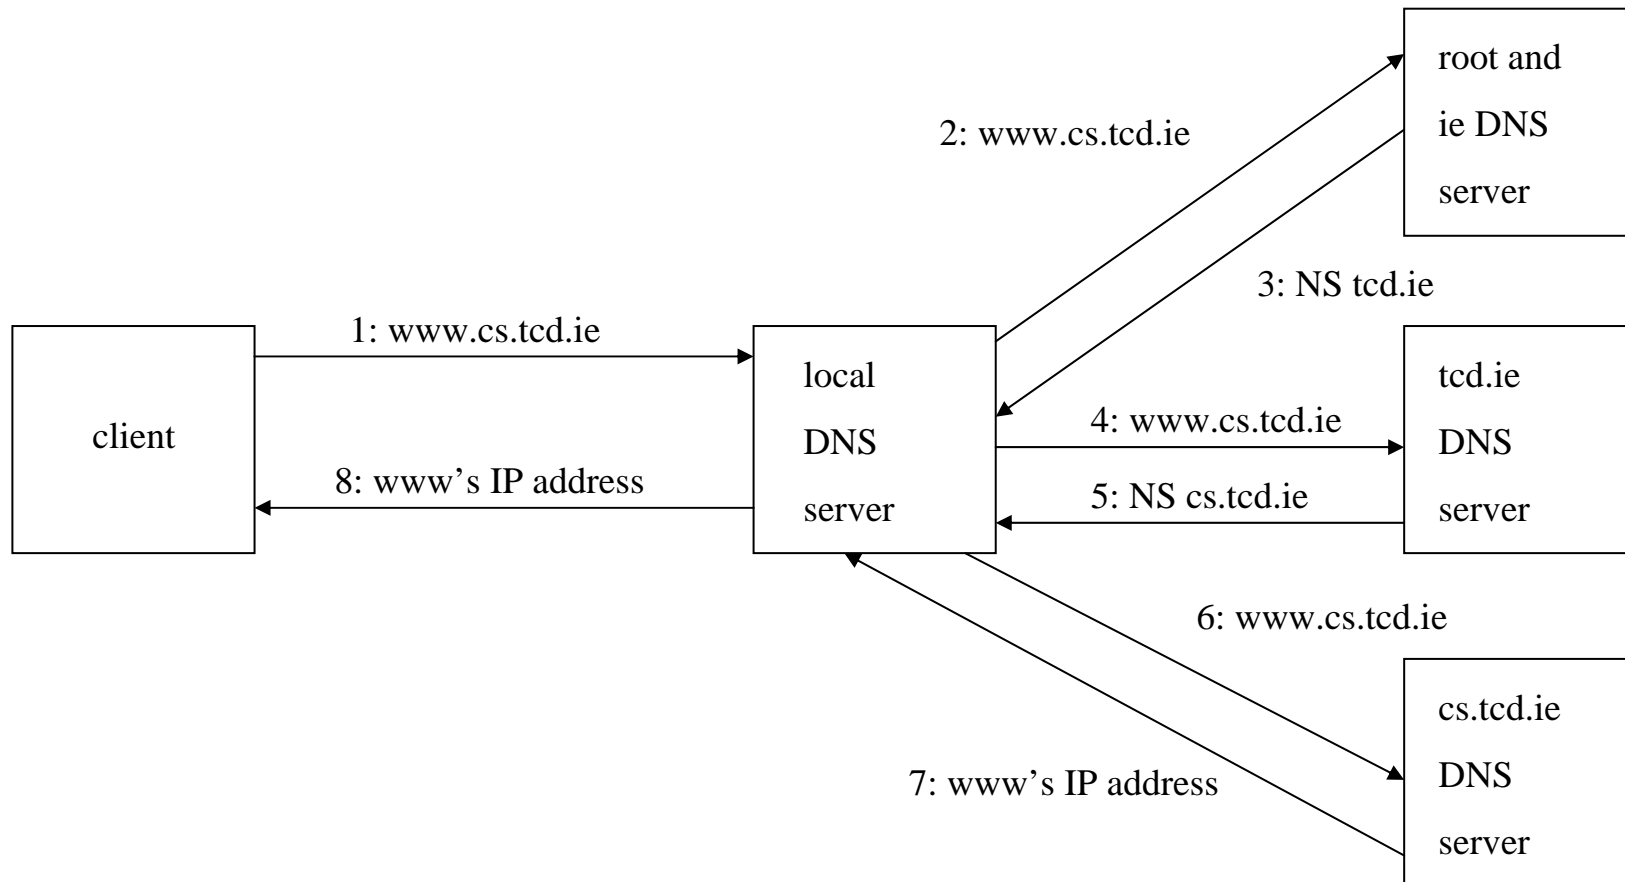
# DNS Protocol and Messages (2/2)



# Inserting Records Into DNS

- Example: Just created startup ``Network Fun''
- Register name networkfun.com at a registrar (e.g., Network Solutions, Inc.)
  - Need to provide registrar with names and IP addresses of your authoritative name server (primary and secondary)
  - Registrar inserts two RRs into the com TLD server:
    - \* (networkfun.com, dns1.networkfun.com, NS)
    - \* (dns1.networkfun.com, 212.212.212.1, A)
- Put in authoritative server (dns1.networkfun.com) Type A record for www.networkfun.com and Type MX record for networkfun.com
- How do people get the IP address of your web site?

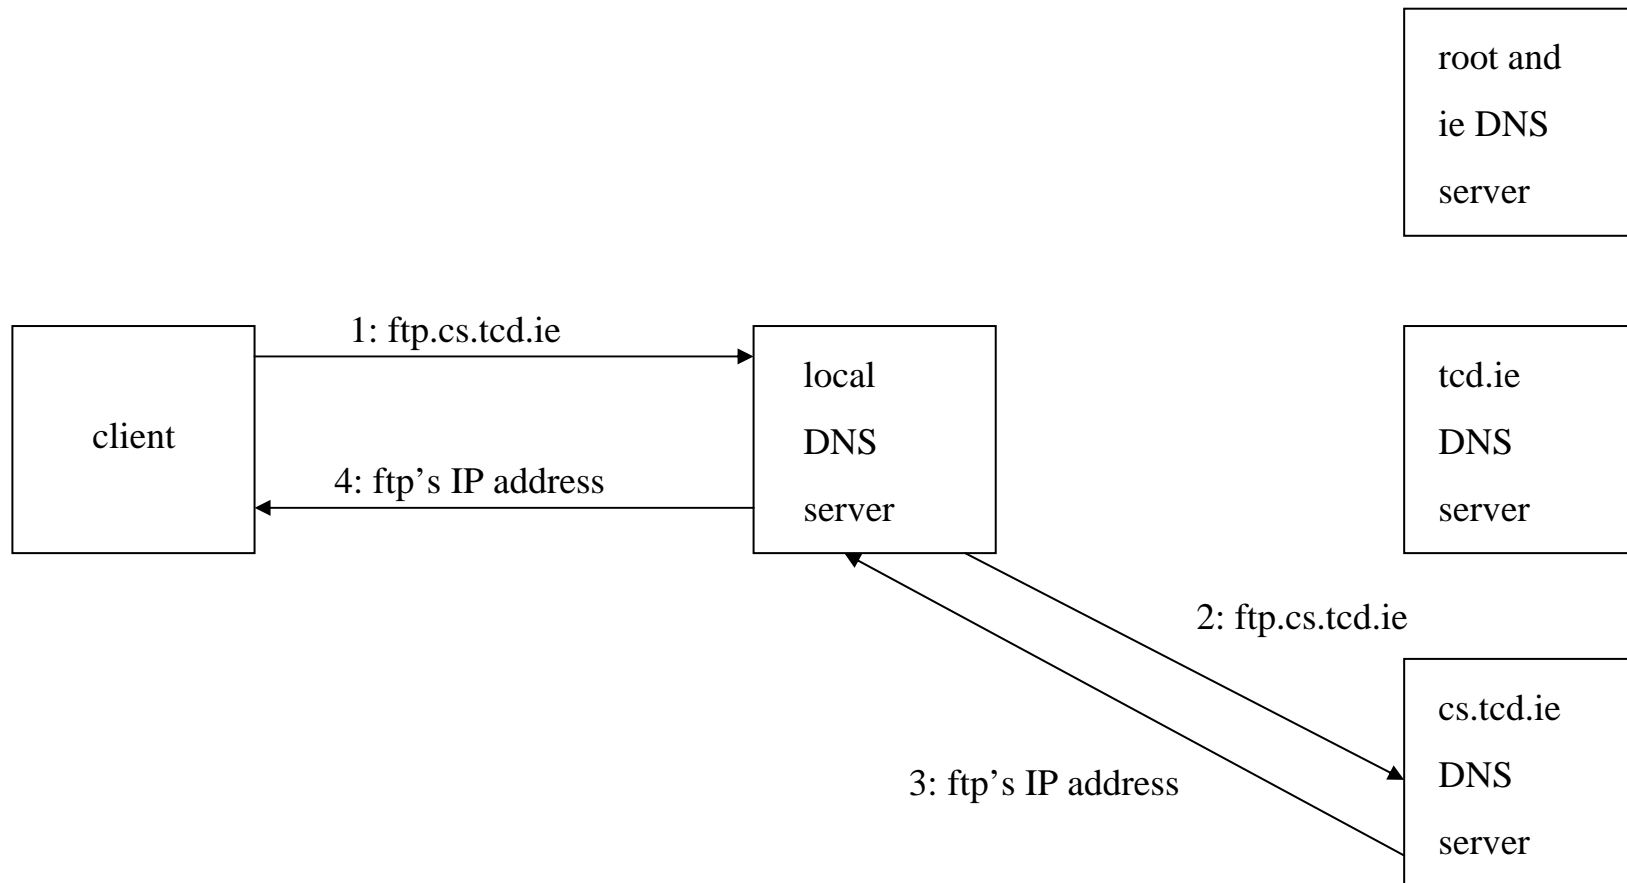
# DNS Lookup Example



# DNS Caching

- DNS responses are cached:
  - Quick response for repeated translations
  - Other queries may reuse some parts of lookup
    - \* NS records for domains
- DNS negative queries are cached
  - Don't have to repeat past mistakes
  - E.g. misspellings
- Cached data periodically times out
  - Lifetime (TTL) of data controlled by owner of data
  - TTL passed along with every record

# Subsequent Lookup Example






# Reverse DNS

- Given numeric IP address, find DNS name
- To find 150.10.20.1:
  - Query 1.20.10.150.in-addr.arpa
  - Get back server.acme.com

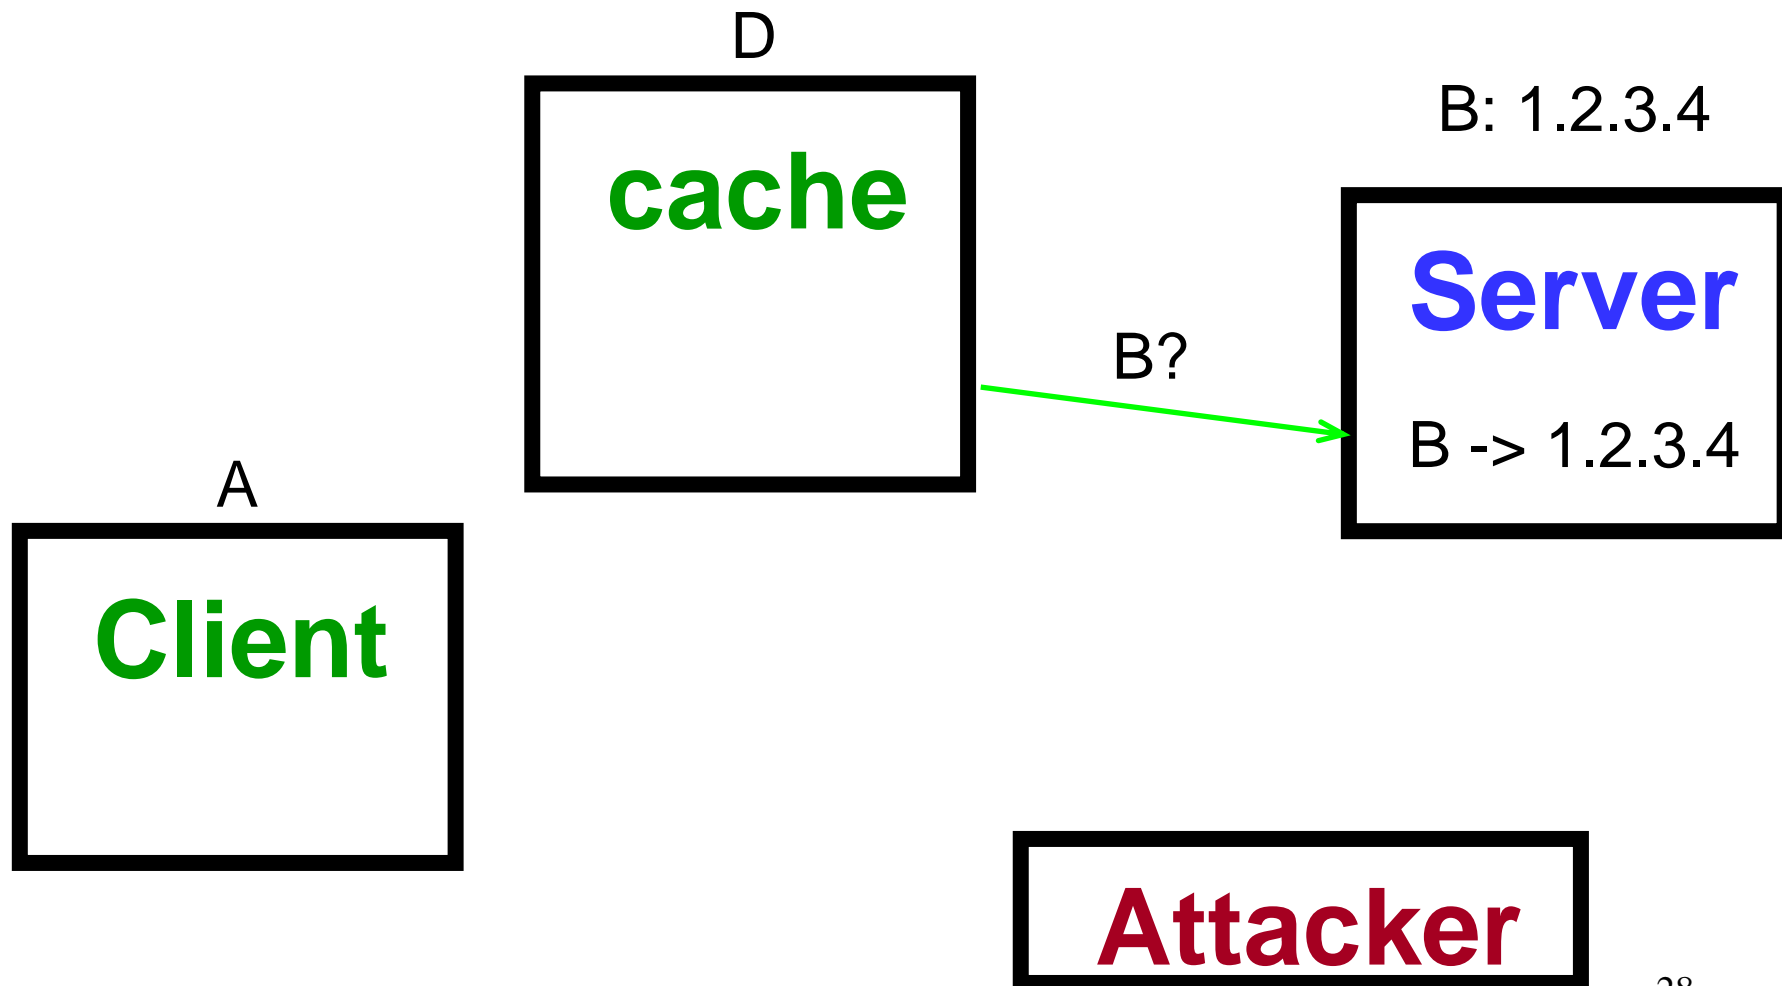
# DNS Attacks

- Cache poisoning
- Reverse DNS attack
  - Known as the Bellovin/Mockapetris attack


# DNS Lookup: A asks D for B's IP address




# DNS Lookup: D asks B (or someone who knows about B)



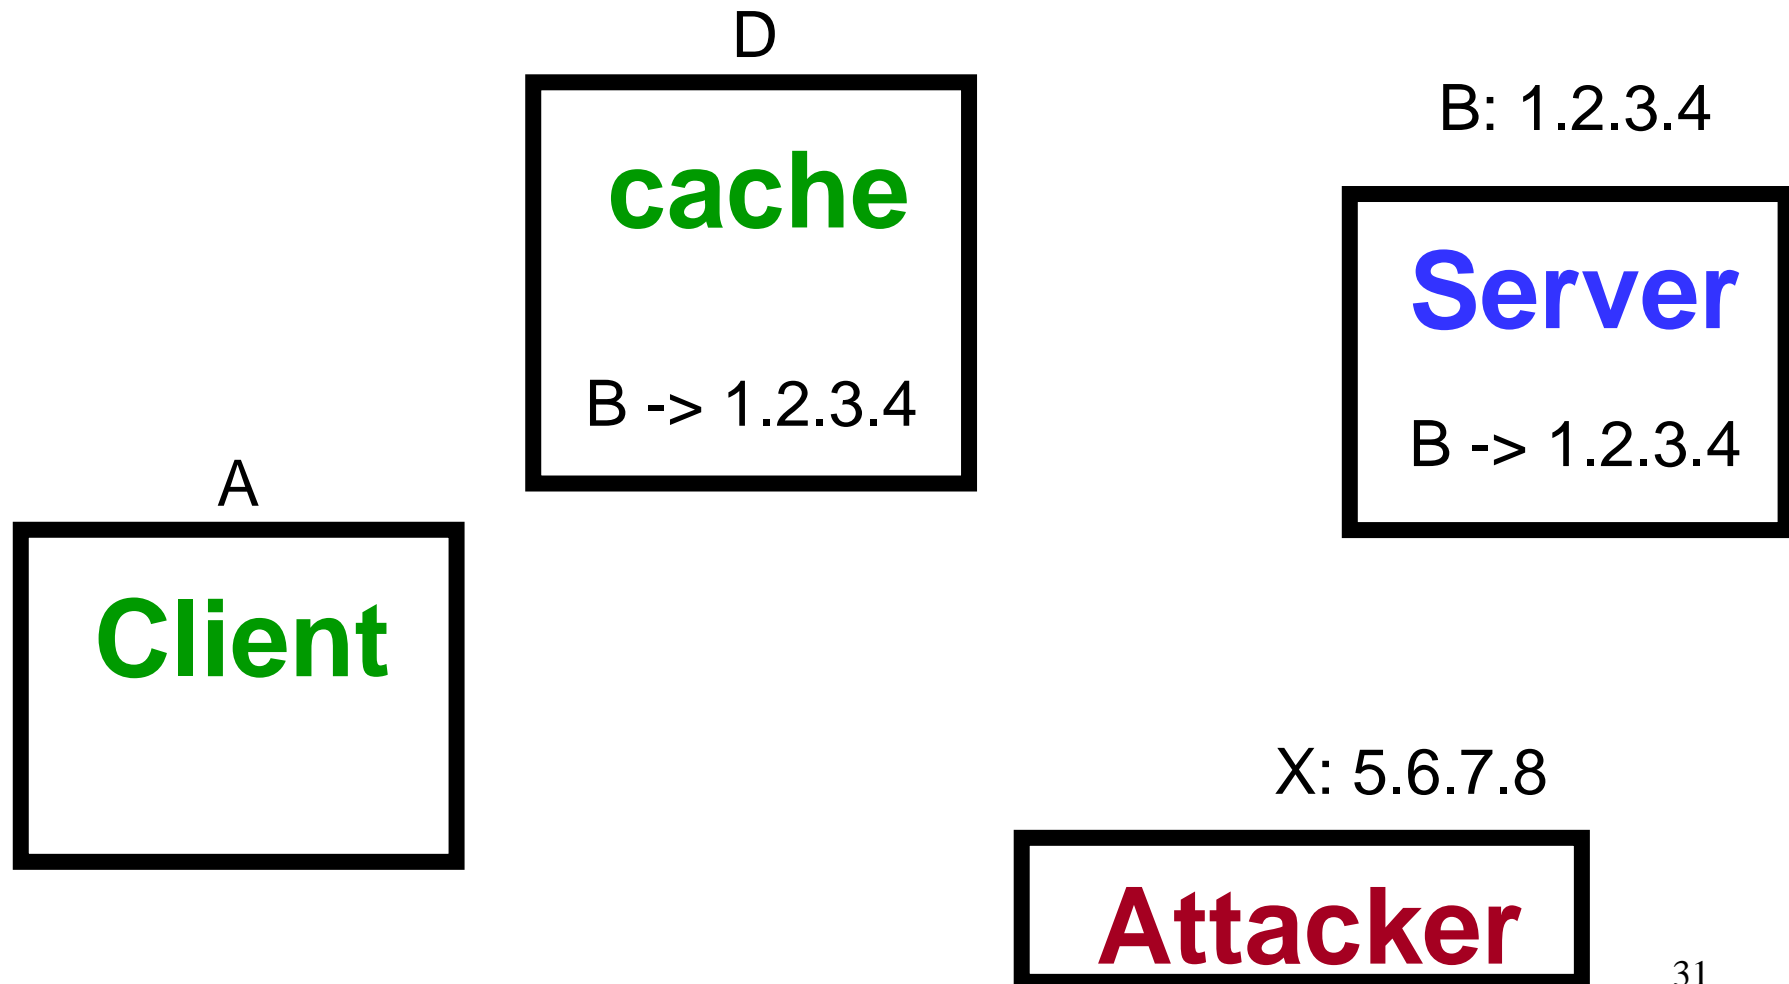
# DNS Lookup: B answers, D caches the answer, and tells A



# DNS Lookup: A uses the answer

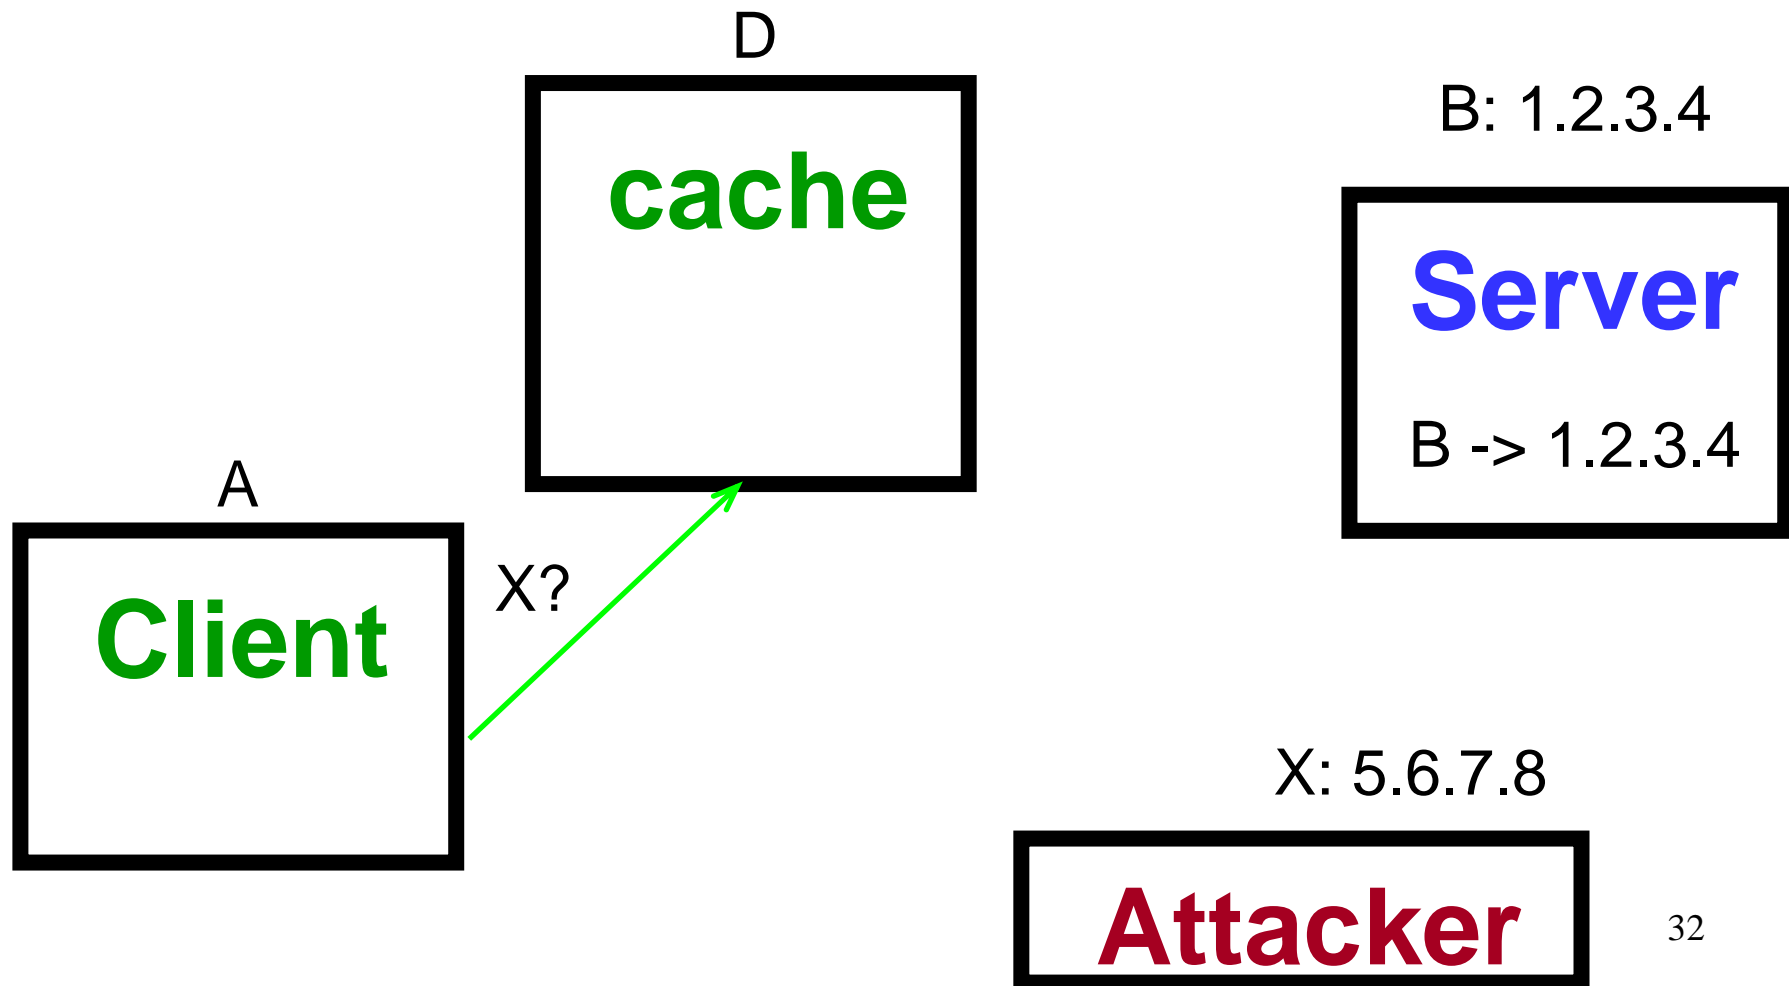


# DNS Lookup: D remembers the answer for a given period



# DNS Cache Poisoning Attack:

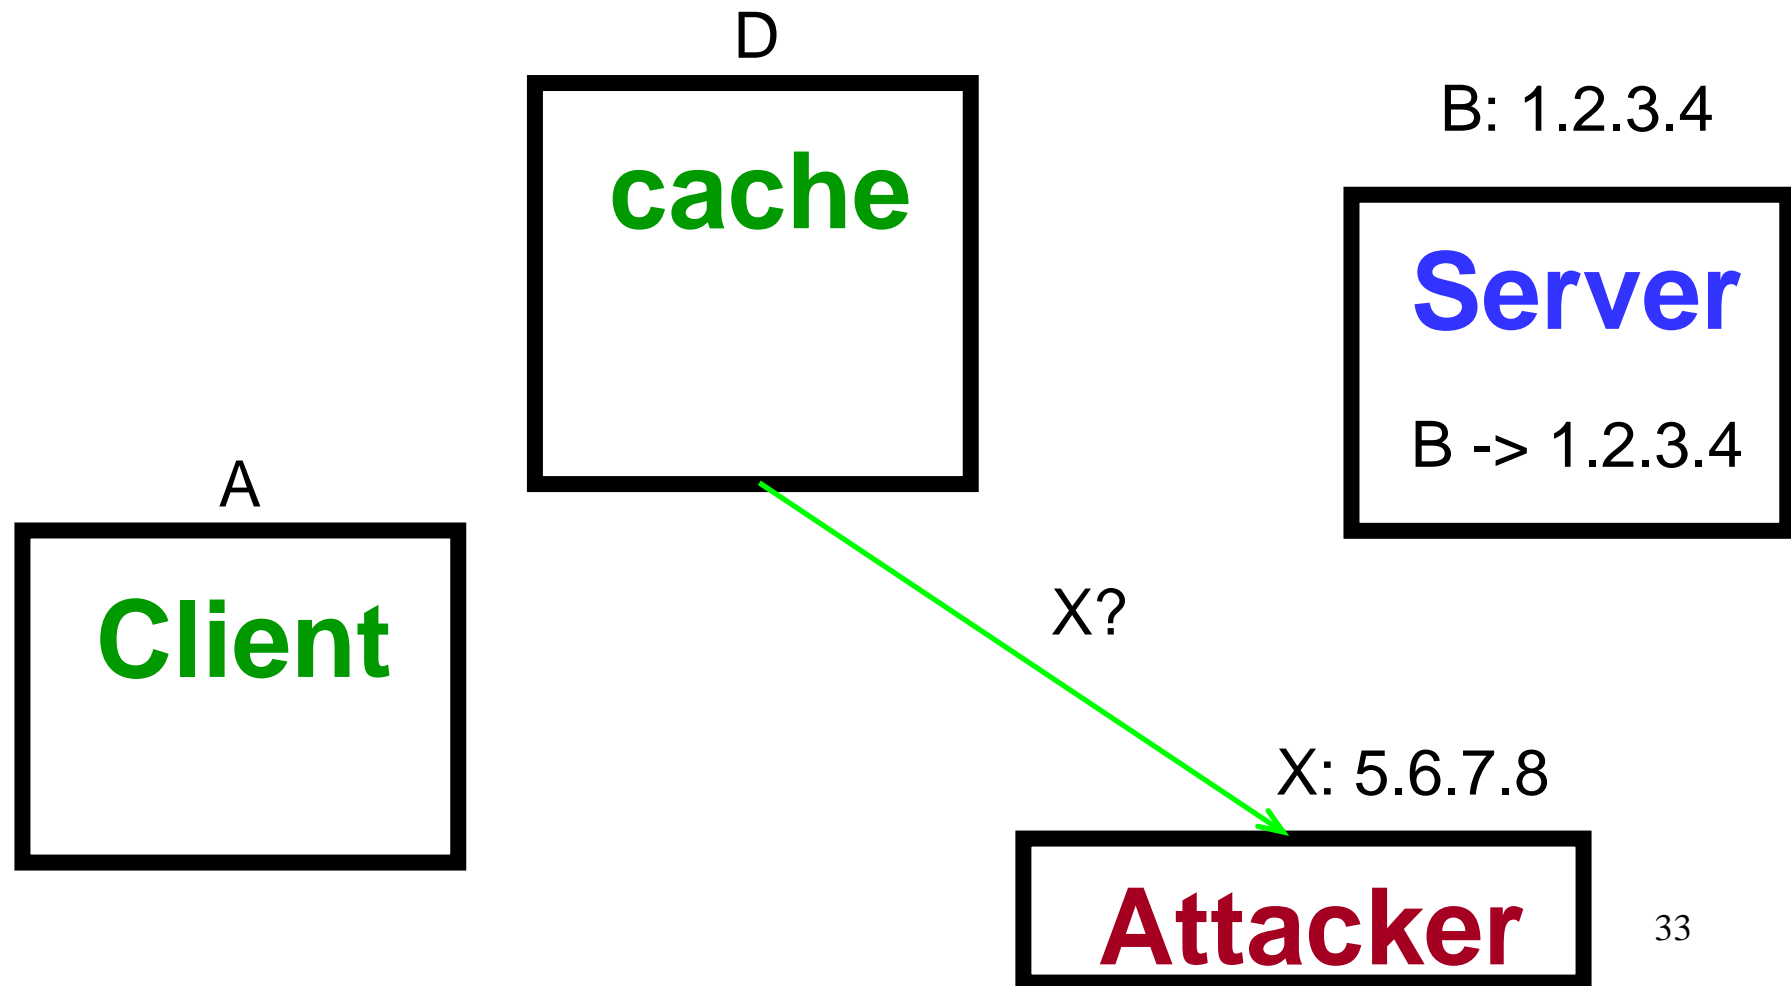
Attacker C arranges for D to ask him  
a question



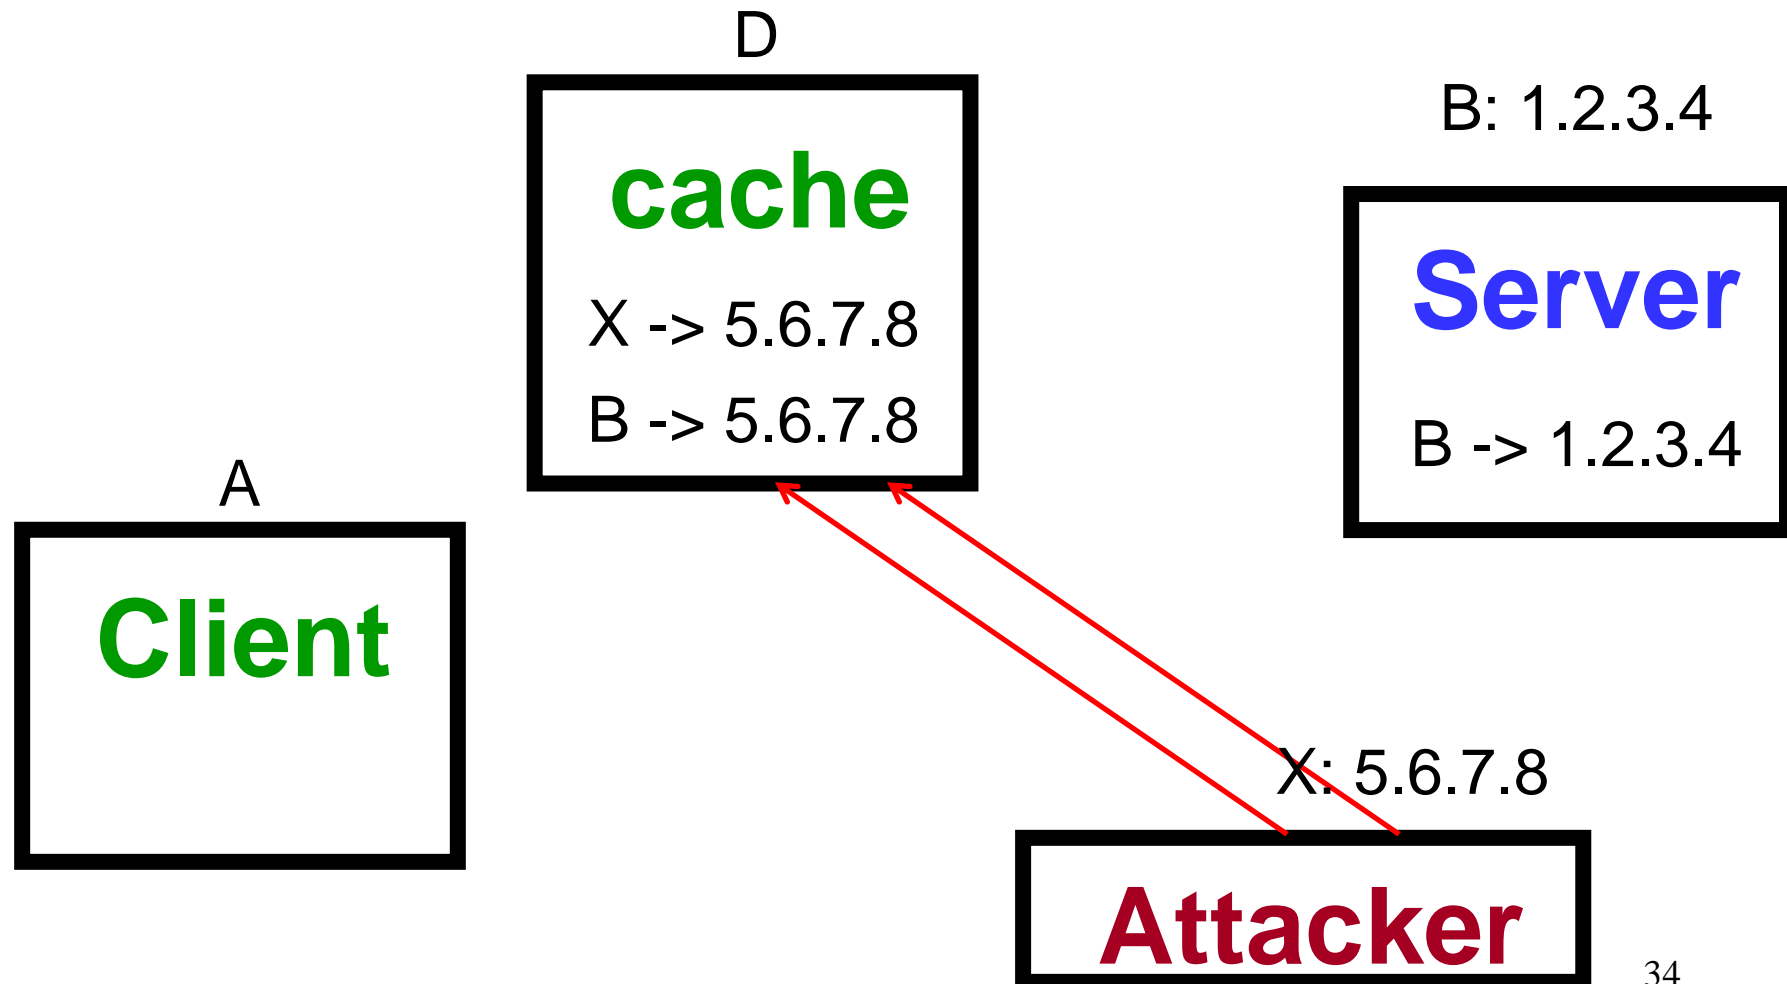


# DNS Cache Poisoning Attack:

Attacker C arranges for D to ask him  
a question

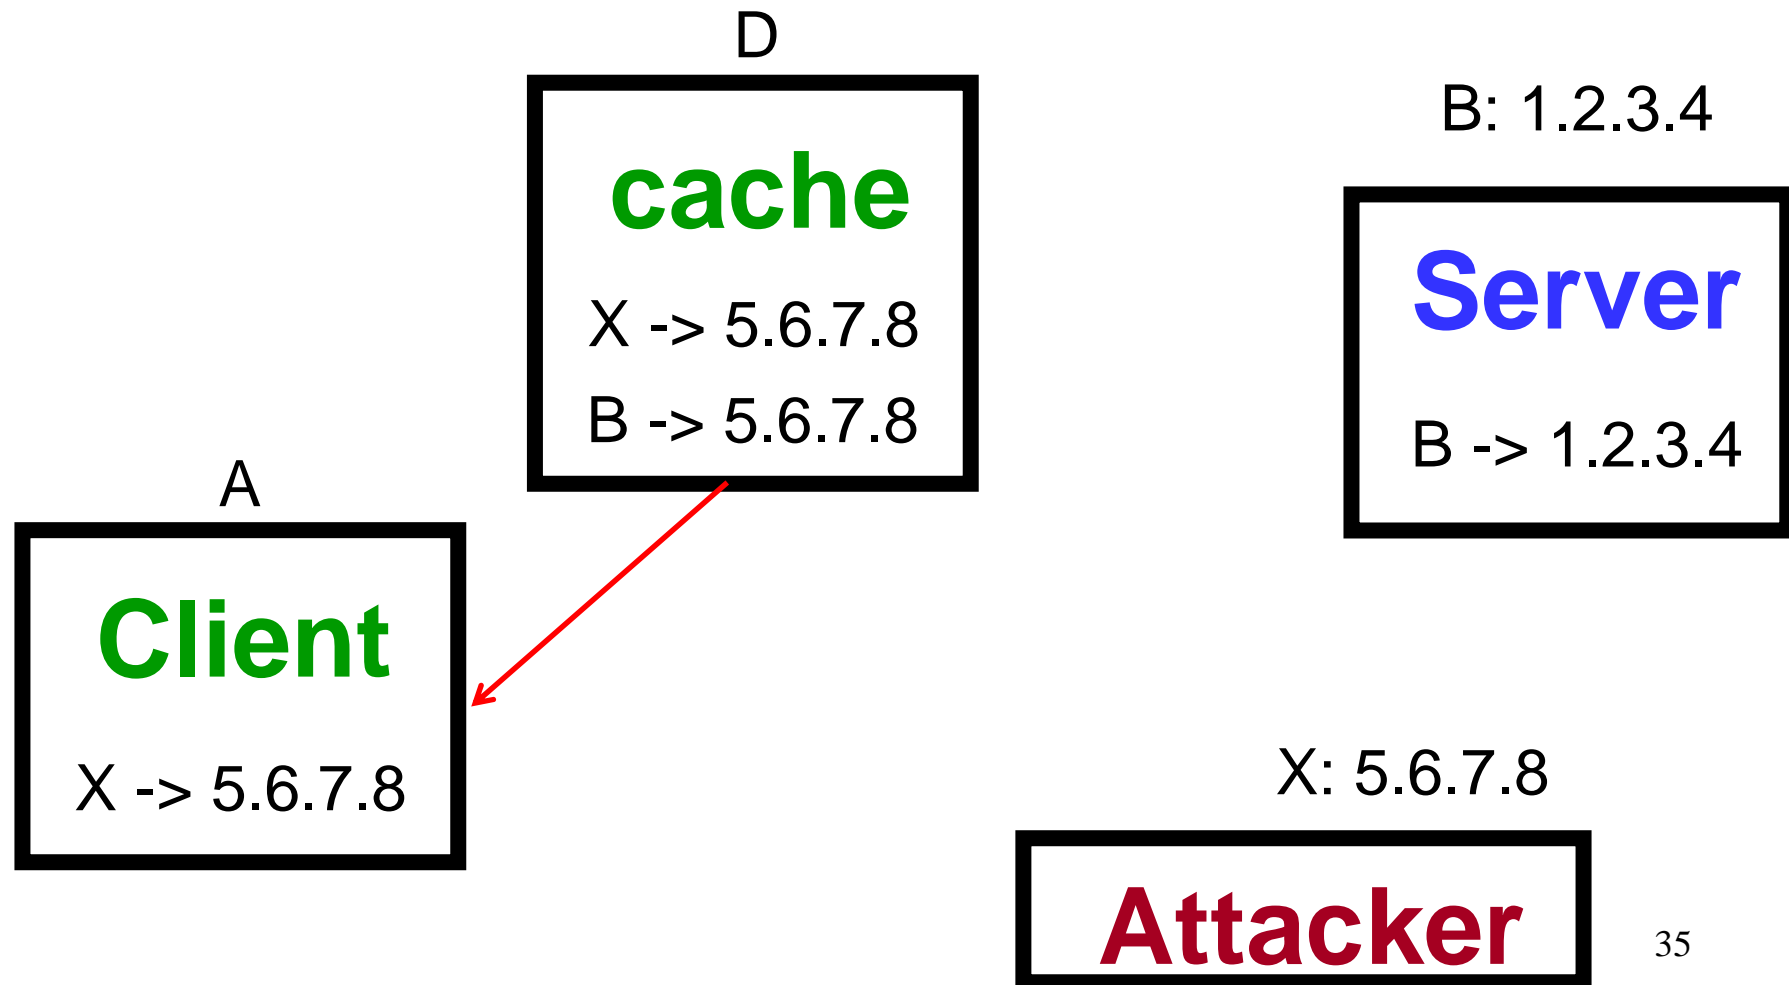


# DNS Cache Poisoning Attack: The attacker gives an answer, plus...



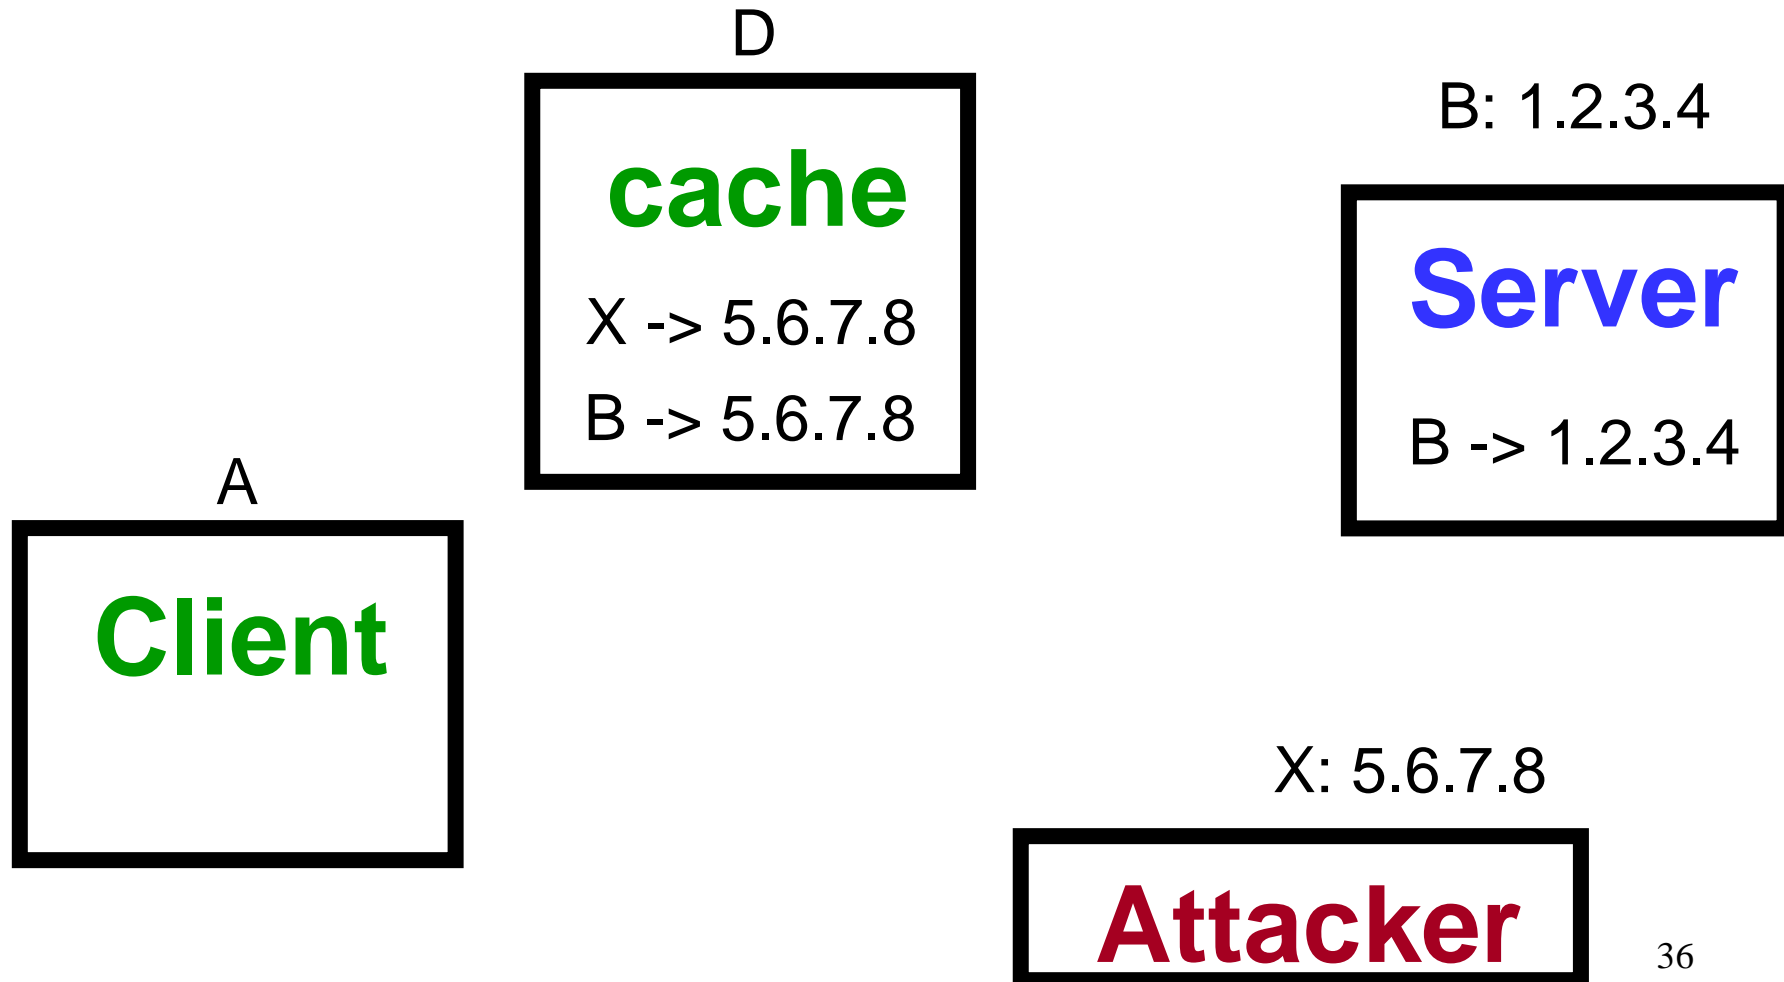
# DNS Cache Poisoning Attack:

A gets his answer, and uses it



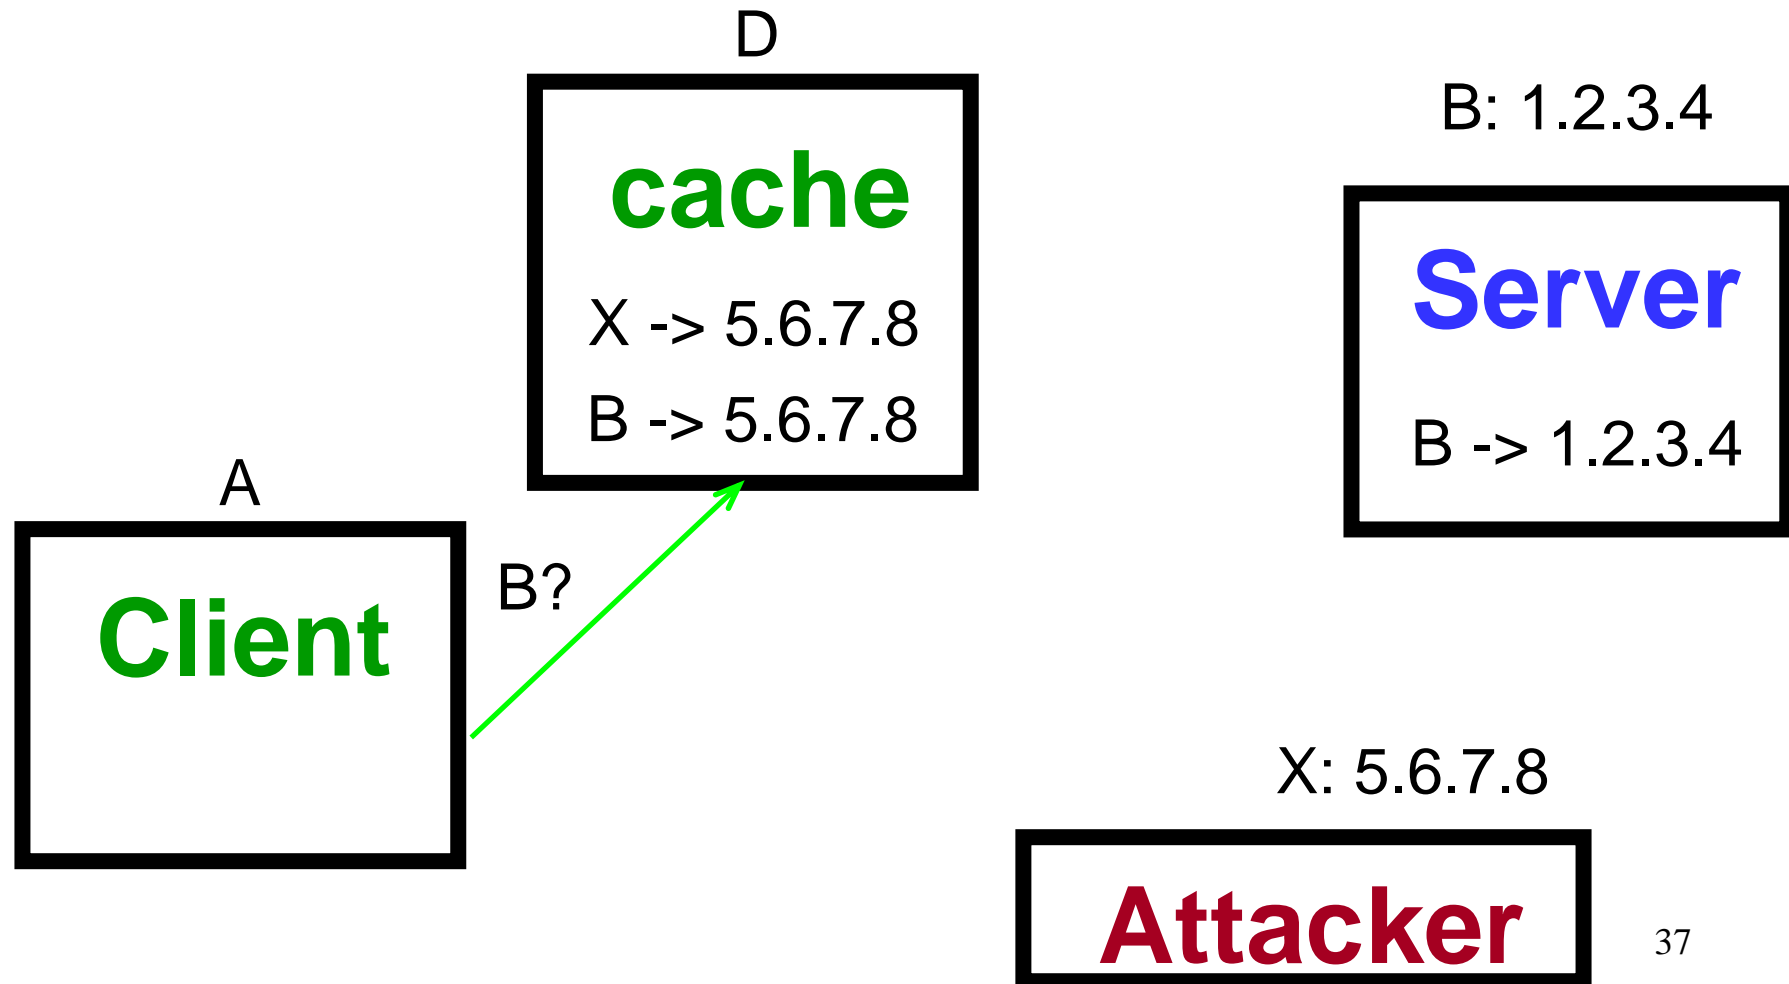
# DNS Cache Poisoning Attack:

The cache has an extra answer

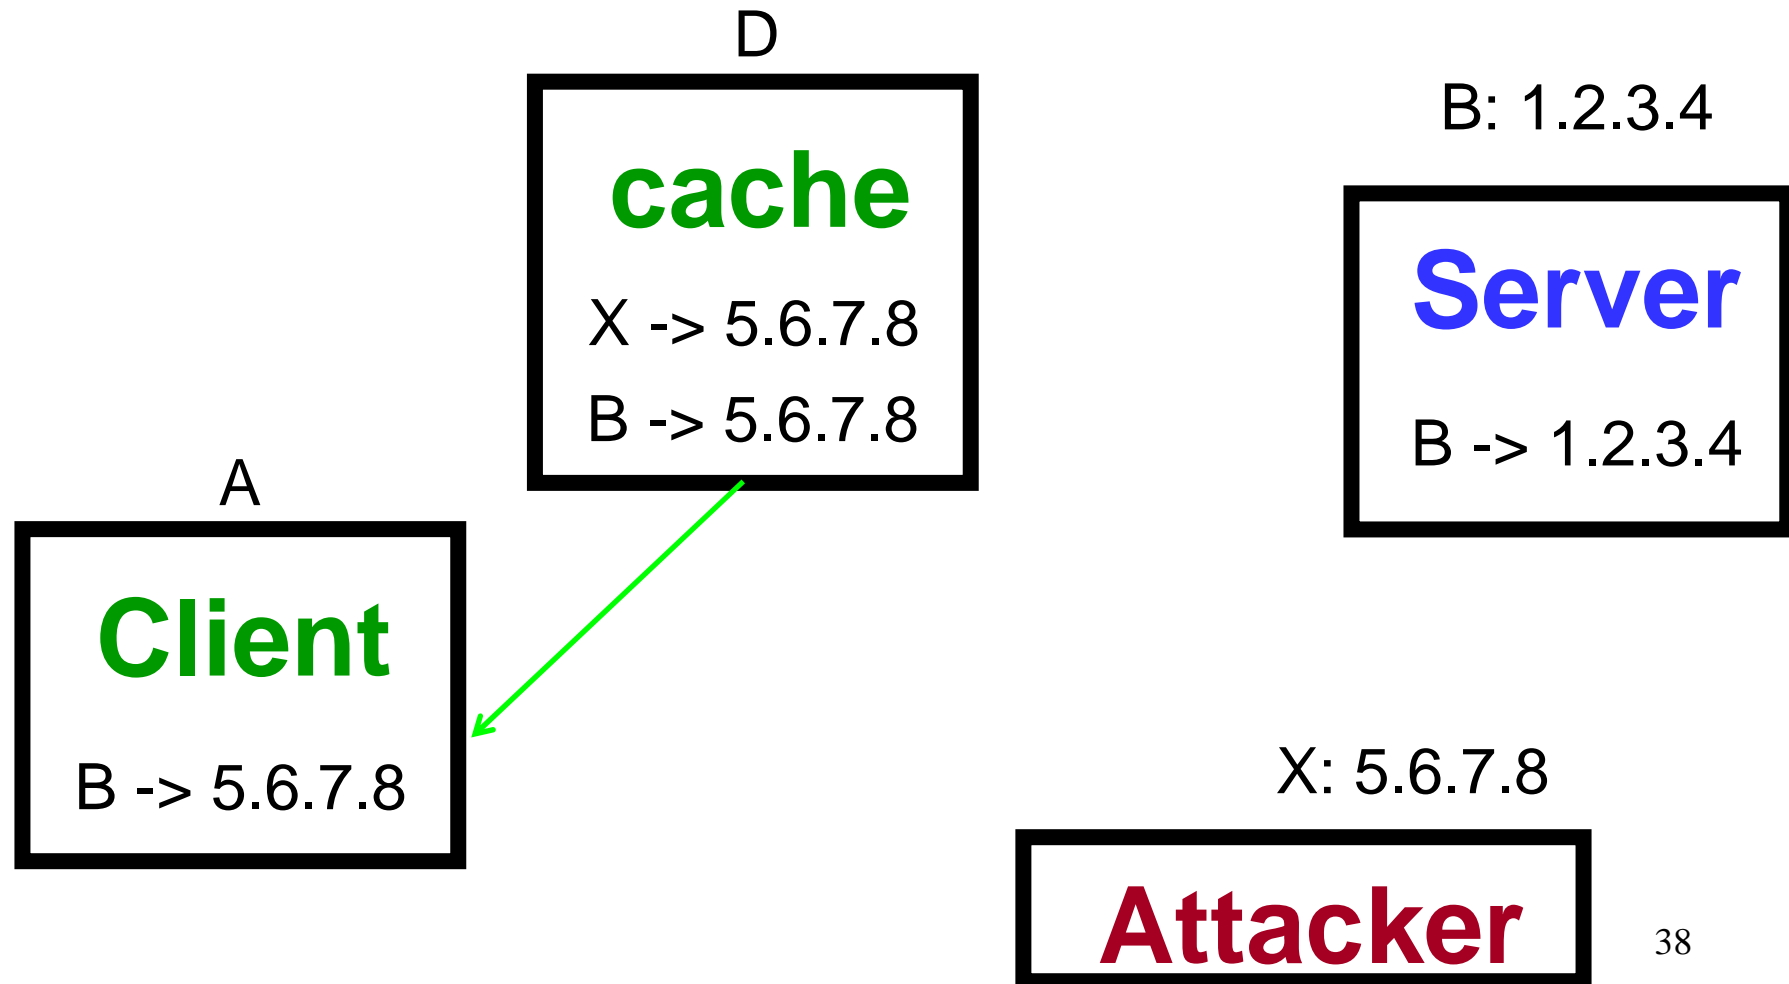


# DNS Cache Poisoning Attack:

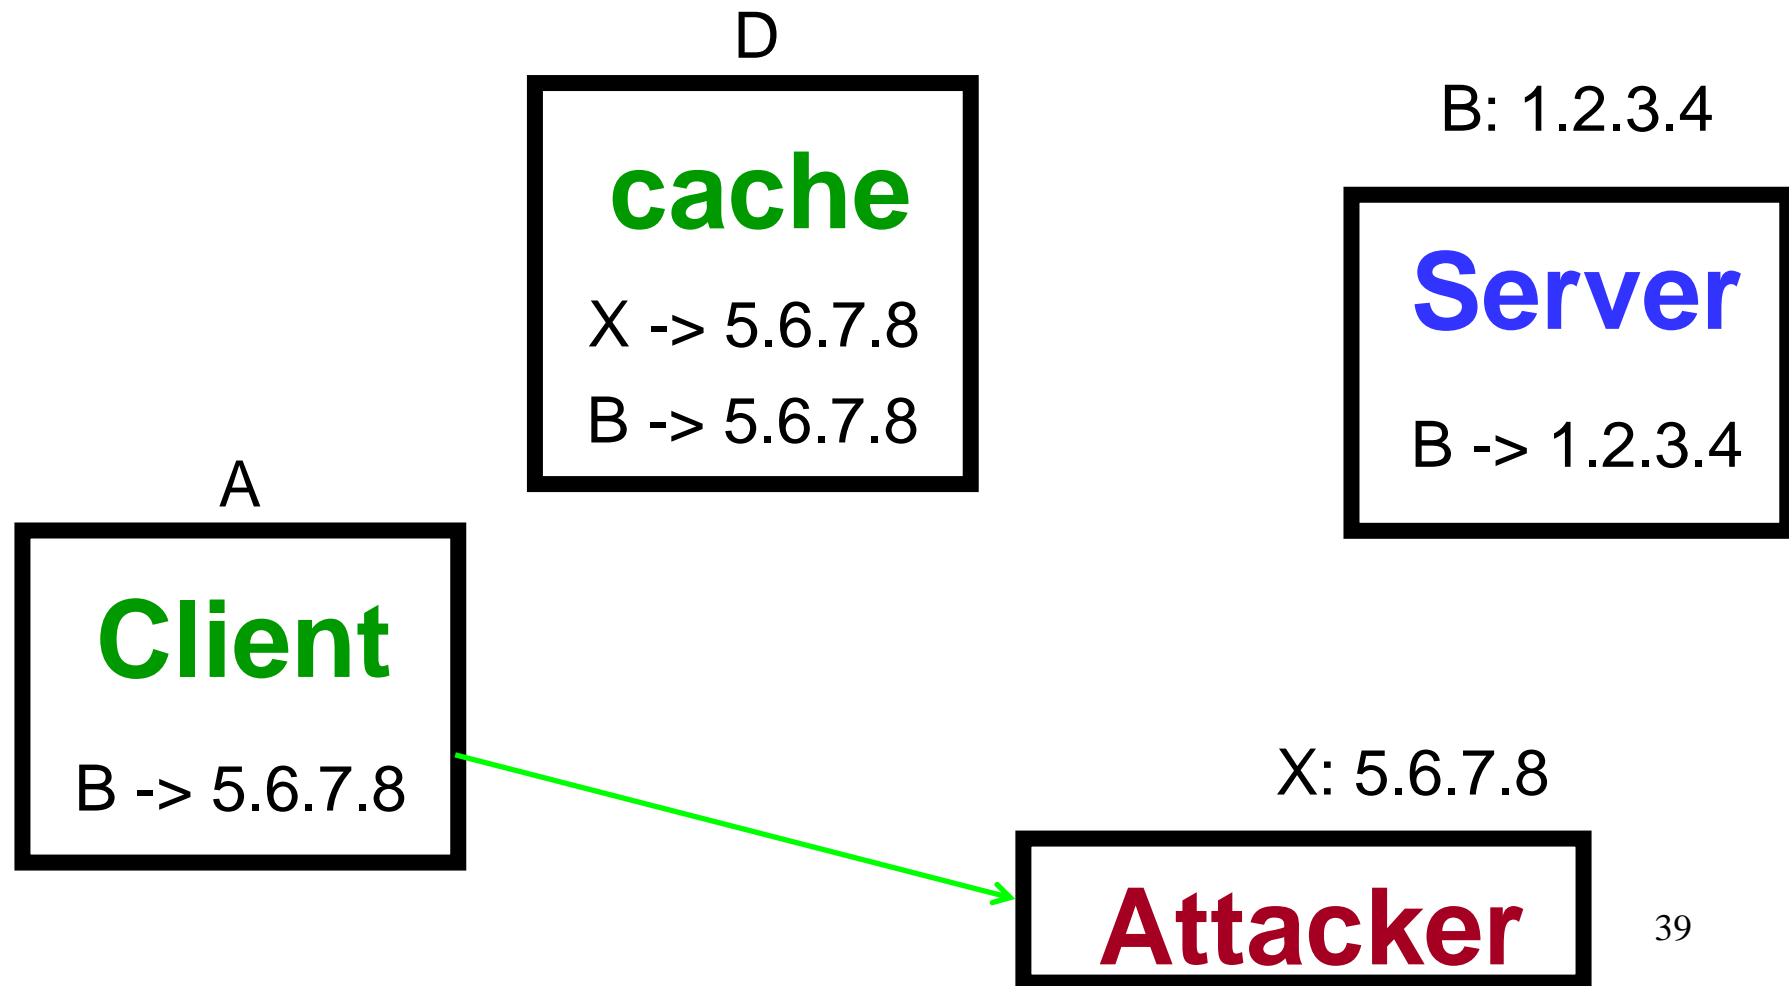
Now A asks for B's address



DNS Cache Poisoning Attack: D  
``knows'' the answer already, and  
returns it



# DNS Cache Poisoning Attack: A uses the answer



# DNS Cache Poisoning

- Older versions of bind fall for this
- You can even send an answer without a query, to some implementations!
- DNS responses can be spoofed to
  - What if the query gets two answers: Use the first?!
- DNSsec fixes this



# Bellovin/Mockapetris Attack

- Trust relationships use DNS names
  - /etc/hosts.equiv contains ntrg.cs.tcd.ie
- Requests come with numeric IP source address
  - Use reverse DNS to find DNS name
  - Decide access based on /etc/hosts.equiv

# Attack

- Gain control of DNS service for domain
- Select target machine in domain
- Find trust relationships
  - SNMP, finger can help find active sessions, etc.
  - Example: Target trusts host1
- Connect:
  - Attempt rlogin from compromised machine
  - Target contacts reverse DNS server with IP addr1
  - Use modified reverse DNS to say addr1 is host1
  - Target allows rlogin

# Defense Against This Attack

- Double-check reverse DNS:
  - Modify rlogind, rshd to query DNS server
  - See if DNS name maps to numeric IP address
- Authenticate entries in DNS tables:
  - DNSsec
  - Requires some form of PKI...