

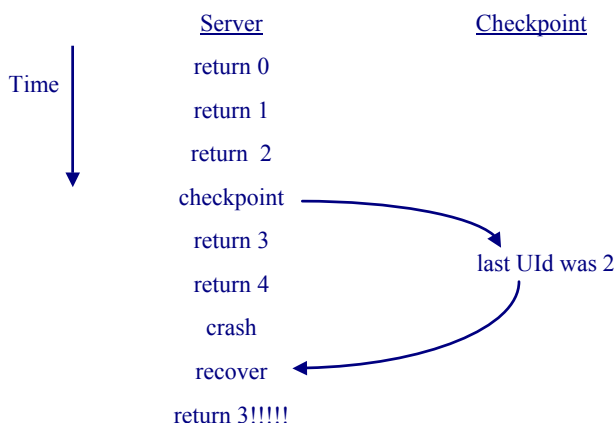
Backward error recovery

- Interested in implementing servers that tolerate machine crashes
 - processor service with partial amnesia crash failures (contents of primary memory are lost)
- *Backward error recovery* refers to the situation where we revert (roll-back) to a previous (saved) state on failure and continue from that point
- Simplest strategy is to periodically *checkpoint* the server's state to disk and restart the server from its last saved checkpoint

For example:

- A server which generates unique identifiers (UIDs)
 - assume that a new UID is obtained by incrementing a counter

Checkpointing 1/4



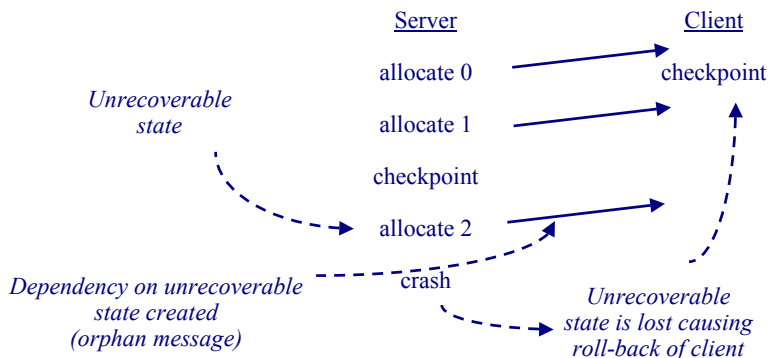
Multiple clients might be allocated the same *unique* identifier - WHY!?

Checkpointing 2/4

- Basic problem is that changes made to the server between the last checkpoint and the crash are *unrecoverable*
 - i.e. its as if UIDs 3 and 4 were never allocated - history has been undone!
- If we consider the server as a state machine, then the states in which UIDs 3 and 4 were allocated have been *lost*
- This is not a problem if the lost states were not visible outside of the server
- However in the example the clients who were originally allocated UIDs 3 and 4 have observed these lost states
- In particular, these clients *depend* on the lost states
 - c.f. isolation property of transactions (dirty read)

Checkpointing 3/4

- One solution is to roll-back any client that has a dependency on a lost state of the server to a state where no such dependency exists

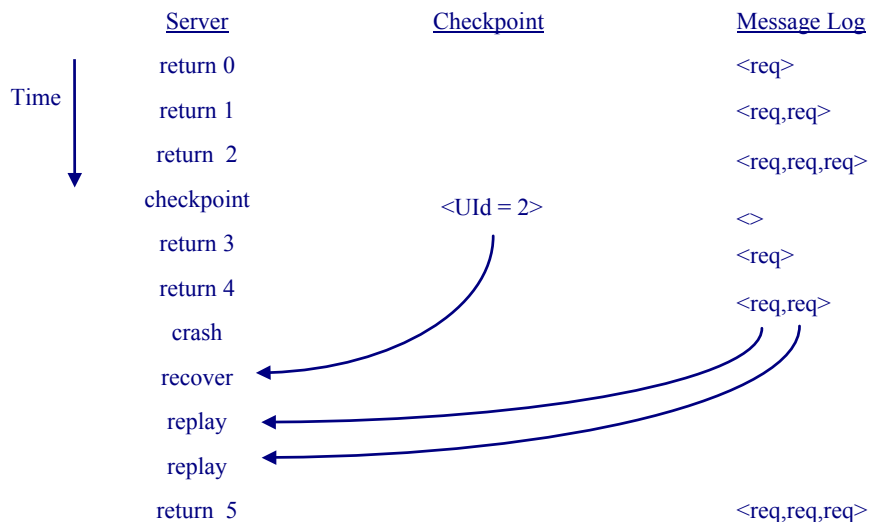


- Of course any process that has a dependency on an unrecoverable state of the client must also roll back - cascading roll-backs (*the domino effect*)

Checkpointing 4/4

- Possible solution to this problem is to *prevent* dependencies on unrecoverable states from arising
- Could checkpoint before replying to each request
 - synchronous disk write - must make sure that the block is flushed to disk before replying
 - may be practical for this example but not if the state to be checkpointed is larger
- Better solution is to keep a *log* of all requests received since last checkpoint and replay these in order after restoring checkpoint
 - again synchronous write - but only of request message

Using checkpoints and message logs 1/2



Using checkpoints and message logs 2/2

- Each request must be logged before result is returned
 - in case crash between returning result and writing log record
 - still pay one disk write per request
- Requires that server execution is *deterministic*
 - so replaying the messages in the log will yield the same state after recovery as before failure
 - needs care if server is multi-threaded!
- Any other inputs to the server must be logged too
 - e.g. terminal input
- Sending of (duplicate) reply messages must be suppressed during replay
 - unless clients can detect duplicates
- Keep reply messages in case clients retransmit requests
 - e.g. crash after logging message but before replying

Sender-based message logging 1/6

See [Johnson and Zwaenepoel 1987]

- A cheaper alternative to logging requests on disk
- Supports recovery from a *single* failure at a time
- Processes must be deterministic
- Basic idea is to keep the message log in the sender's volatile memory
- On recovery each process restores its state from its checkpoint and asks its correspondents to resend their messages to it
- Messages must be replayed in the same order as originally processed
- Each message carries a Send Sequence Number (SSN) for duplicate suppression

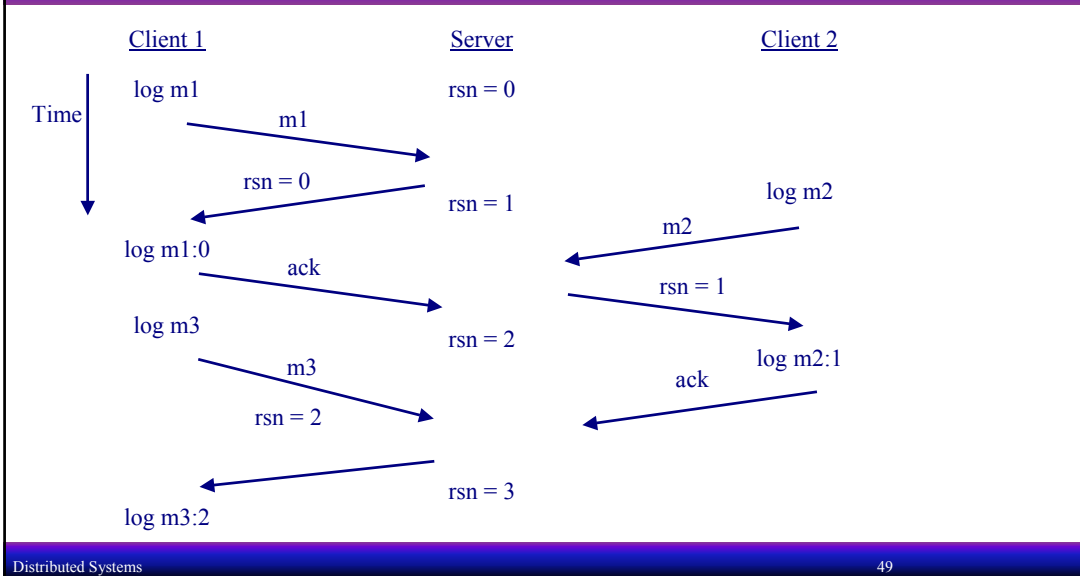
Sender-based message logging 2/6

- To solve the ordering problem each receiver assigns a Receive Sequence Number (RSN) to each message that it receives
- Messages are replayed in RSN order
- RSNs are returned to senders and must be added to the corresponding messages in the log
- A message for which no RSN has yet been logged is known as a *partially logged* message

Sender-based message logging 3/6

- The state resulting from processing a message is unrecoverable while the message is partially logged
- The receiver may not send further messages or perform I/O while there are any partially logged messages
- Sender must ack each RSN
- During recovery partially logged messages may be replayed in any order after all fully logged messages
 - e.g. if the server fails without returning an RSN for some message

Sender-based message logging 4/6



Sender-based message logging 5/6

During recovery:

- process is restored from checkpoint (includes its highest used SSN/RSN)
- process broadcasts to solicit messages to it with RSNs greater than the checkpoint RSN (or no RSN)
- fully-logged messages are processed in RSN order
- if recovering process (re)sends any messages these will have the same SSN as before and are added to its local log
- suppression of these messages is the responsibility of their receivers
- the receiver of such a duplicate must return its RSN or an indication that the message need not be logged (if it has checkpointed)
- partially logged and new messages now processed in any order

Sender-based message logging 6/6

- Checkpoint includes state of process, its current message log and its highest used SSN and RSN
- Once a process checkpoints, log records for messages sent to it with RSNs lower than the checkpoint RSN can be discarded


Optimistic approaches to checkpointing

- Both the basic and sender-based checkpointing/message logging schemes *prevent* dependencies on unrecoverable states from arising
- These protocols are *pessimistic* in that they assume that failures will occur and prevent communication/output that could lead to inconsistency
- Possible to have an *optimistic* protocol which will allow communication between processes to proceed even while some messages are not logged
- Such protocols are optimistic in that they assume that logging will normally complete without failure
- Recovery in this case may involve rolling back multiple processes
- May support *non-deterministic* processes
- Still necessary to delay output from unrecoverable states

Distributed checkpointing 1/3

- Yet another approach is to take a *distributed checkpoint* covering a set of communicating processes rather than a single process
- Such a checkpoint must capture a *consistent* state of the processes
- The major criterion for consistency is that after recovery no process should be in state where it has received an *orphan* message
- Since there is no message log, such protocols must handle message loss due to processor crashes
- Like optimistic approaches, distributed checkpointing may support *non-deterministic* processes
- Must not allow output from unrecoverable states

Distributed checkpointing 2/3



- Note that C_1 represents a consistent checkpoint while C_2 does not
- In C_2 , p_2 has received an orphan message (m_5) from p_3
- Also note that rolling back to C_1 may cause message loss (m_1 , m_3)

Distributed checkpointing 3/3

- Design algorithms for taking a consistent checkpoint of the state of a set of communicating processes and for rolling back to the checkpoint state after the failure of one of the processes
- For extra brownie points, ensure that your algorithms only require that the minimum number of processes checkpoint/rollback as required
- Hints:
 - Taking a checkpoint may be initiated by any process
 - Your algorithm may give rise to lost messages

Taking a distributed checkpoint 1/9

See [Koo and Toueg 1987]

- Assume that we can identify a set of mutually consistent local states of the required processes
- Failures during checkpointing might prevent one or more processes from checkpointing their local states
- So, we need to ensure that all required processes checkpoint or none do
 - taking a checkpoint should be an atomic (all-or-nothing) operation
 - even in the presence of failures
- Algorithm requires two rounds of communication initiated by process that decides to take the checkpoint - *the initiator*
 - essentially 2PC

Taking a distributed checkpoint 2/9

Round one:

- Initiator
 - takes a *tentative* checkpoint
 - requests all required processes to make tentative checkpoints
- Other processes
 - may/may not take tentative checkpoints
 - inform initiator of their action
 - may *not* send any further messages until algorithm is complete

Round two:

- If initiator establishes that all processes have taken tentative checkpoints
 - it makes its checkpoint *permanent*
 - requests others to do likewise
- Other processes comply with initiator's request

Taking a distributed checkpoint 3/9

Why no orphan messages?

- Suppose p_1 sends m_1 to p_2
- For m_1 to be an orphan
 - reception of m_1 would have to occur before p_2 takes its checkpoint
 - sending of m_1 would have to occur after checkpoint of p_1
- No sending allowed after taking checkpoint until all other processes have checkpointed!


Taking a distributed checkpoint 4/9

What about failure?

- If process is unable to reply in round one, its reply is assumed to be negative and no checkpoint is taken
- Once process has replied in round one, it must follow initiator's decision
- If it fails it may need to contact initiator to determine outcome
- If the initiator fails before making checkpoint permanent, no checkpoint is taken
- If the initiator fails after making checkpoint permanent, all processes will make checkpoints permanent

Taking a distributed checkpoint 5/9

- Note that it may not be necessary for every process to checkpoint



- If a checkpoint was taken at C_1 and p_1 decides to initiate a checkpoint after receiving m_4 , only p_1 and p_2 need to take new local checkpoints to have a consistent distributed checkpoint C_2

Taking a distributed checkpoint 6/9

- So, every message carries a send sequence number (SSN)
- Each process r , keeps track of the SSN of the last message it received from every other process since its last checkpoint


$$\text{last_recd}_r(s) \text{ for all } s \neq r$$
- Each process s , keeps track of the SSN of the first message that it sent to every other process since its last checkpoint

$$\text{first_sent}_s(r) \text{ for all } r \neq s$$
- When p_i asks p_j to checkpoint, it includes $\text{last_recd}_{p_i}(p_j)$ in the request
- p_j needs to checkpoint only if

$$\text{last_recd}_{p_i}(p_j) > \text{first_sent}_{p_j}(p_i)$$
- Each process p_i keeps a list of the processes from which it has received messages since its last checkpoint - cohorts_{p_i}


Taking a distributed checkpoint 7/9

- Checkpoint request is only sent to cohorts of initiator
- May be propagated to *their* cohorts transitively



Taking a distributed checkpoint 8/9

- Say p_2 initiates a checkpoint after receiving m_5 in the following scenario:



Taking a distributed checkpoint 9/9

- On recovery of a failed process, could rollback every process in the system to its last permanent checkpoint
 - as an atomic action requiring two rounds of communication
- If the failed process p_i hasn't communicated to p_j it may not be necessary to rollback p_j
- Each process s , keeps track of the SSN of the last message that it sent to every other process before it took its last permanent checkpoint
 - $\text{last_sent}_s(r)$ for all $r \neq s$
- When p_i asks p_j to rollback, it includes $\text{last_sent}_{p_i}(p_j)$ in the request
- p_j needs to rollback only if

$$\text{last_recd}_{p_j}(p_i) > \text{last_sent}_{p_i}(p_j)$$

References

- [Johnson and Zwaenepoel 1987] David Johnson and Willy Zwaenepoel, *Sender-Based Message Logging*, Proceedings of the 17th IEEE International Symposium on Fault-Tolerant Computing Systems, pages 14-19, 1987.
- [Koo and Toueg 1987] Richard Koo and Sam Toueg, *Checkpointing and Rollback Recovery for Distributed Systems*, IEEE Transactions on Software Engineering, SE-13(1):23-31, January 1987.

See also:

- [Jalote 1994] Pankaj Jalote, *Fault Tolerance in Distributed Systems*, Prentice Hall, 1994, ISBN 0-13-301367-7, Chapter 5.

Process replication

- Another approach to providing servers that can tolerate machine crashes is to replicate server processes

Basically two approaches:

- Passive/loosely-synchronised/primary-backup replication
 - primary process handles each request and replies to clients while passive/backup process collects recovery information
 - backup takes over on failure
 - can have multiple backups
- Active/closely-synchronised replication
 - multiple processes handle each request in parallel

Passive replication

- Each server has a *primary* that accepts and responds to requests and...
- ...a *backup* that takes over when the primary crashes
- An alternative to storing checkpoints/message logs on disk
- Backup has checkpoint and message log
- Trade-off cost of logging against cost of sending requests to both primary and backup processes


TARGON/32 1/4

- TARGON/32 is a fault-tolerant version of UNIX [Borg et al. 1989]
- Goals:
 - recover from any single crash failure
 - ◆ partial amnesia crashes as usual
 - complete transparency (programs do not have to be modified)
 - all processors available to do useful work in the absence of failure
 - tradeoff low overhead during normal execution vs. longer recovery time
- Requirements:
 - processes must be deterministic
 - crashed processes checkpoint must be available to backup
 - all messages since checkpoint must be available to backup and must be replayed in the same order

TARGON/32 2/4

- To have messages replayed in the *same order*, messages from all senders must be *received* by both the primary and backup in the *same order*
- Messages to a server are *multicast* using a *totally-ordered, atomic multicast* to the primary and the backup (and the sender's backup)
 - *atomic* means that multicasts are received by all destinations or none
 - *totally-ordered* means that multicasts are received in the same order at all overlapping destinations
- Messages are copied to the sender's backup so that it knows which messages were sent by primary and can suppress them during recovery

TARGON/32 3/4



Messages m_1 and m_2 are received in the same order at $\{Server_p, Server_b\}$
Log contains messages received by the primary
Discard counter counts messages sent by the primary

TARGON/32 4/4

- Multicast implemented using hardware support
- Checkpoint taken after n messages processed or specified time elapses
- Checkpointing mechanism uses an external pager
 - primary flushes dirty pages (and other state) during checkpoint
 - backup faults most recent checkpoint during recovery
- Replication is optional and a number of different strategies are supported
 - “quarter-backs” - backed up until a crash occurs
 - “half-backs” - new backup created when primary machine rebooted
 - “full-backs” - new backup created asap after crash (not implemented)
- Details are complex (especially supporting UNIX semantics)
 - e.g. signal handling gives rise to non-determinism - checkpoint before handling signal
 - also must be careful about changes to the file system

Active replication

- Requests processed by multiple processes in parallel
- Each request is sent to all replicas
- Requests must be processed in the same order at each process
- Client can collect one or more replies as required


Notes:

- Additional load on system due to replicated processing
- Minimise recovery time after failure - no rollback/recovery/replay
 - suitable for applications with guaranteed response time requirements
- Potentially mask response/arbitrary as well as crash failures by collecting and comparing replies (voting)

Tolerating Byzantine failures 1/14


(The Byzantine Generals Problem)

- Three generals need to agree whether or not to attack an enemy city
- They communicate by exchanging messages (reliably)
- Some of the generals may be traitors
- The traitors may lie but can't impersonate others
- Want to ensure that:
 - all loyal generals decide upon the same plan of action
 - traitors cannot cause the loyal generals to adopt a bad plan




The Byzantine Generals Problem 2/14

- There is no solution!
 - One traitor can subvert the efforts of two loyal generals to reach agreement



The Byzantine Generals Problem 3/14

- *Four* generals need to agree whether or not to attack an enemy city
- They communicate by exchanging messages (reliably)
- Some of the generals may be traitors
- The traitors may lie but can't impersonate others
- Want to ensure that:
 - all loyal generals decide upon the same plan of action
 - traitors cannot cause the loyal generals to adopt a bad plan



The Byzantine Generals Problem 4/14

See [Lamport et al. 1982]

- Given g generals of which at most t are traitors, it can be shown that consensus is possible iff
$$g \geq 3t + 1$$
- No solution for $g \leq 3$
- $g = 4$ generals can tolerate $t = 1$ traitor
- Put another way, to tolerate m faults, we need at least $3m + 1$ processes!
 - if faults are inconsistent and arbitrary

The Byzantine Generals Problem 5/14

- Let v_i be the value for general i and assuming that all generals use the same method for making a decision based on the information they have
- Every loyal general must obtain the same information v_1, \dots, v_n
- If the i^{th} general is loyal, then the value that he sends must be used by every loyal general as the value of v_i

Put another way:

- Any two loyal generals use the same value for v_i
- If the i^{th} general is loyal, then the value that he sends must be used by every loyal general as the value of v_i

The Byzantine Generals Problem 6/14

A general must send an order to his $n-1$ lieutenants such that:

- IC1: All loyal lieutenants obey the same order
- IC2: If the general is loyal, then every loyal lieutenant obeys the order that he sends
- Known as the *interactive consistency conditions*
- No lieutenant can trust an order received from any other participant
- Orders received by other lieutenants are needed to verify the original order
- Lieutenants must *exchange* the orders that they receive

The Byzantine Generals Problem 7/14

- The *oral message* algorithm depends on the number of traitors
 - $OM(t)$
 - assume that a loyal lieutenant follows the algorithm correctly

Assumptions:

- A1: Every message that is sent is delivered correctly
- A2: The receiver of a message knows who sent it
- A3: The absence of a message can be detected
- Algorithm proceeds in *rounds* requiring a *synchronous* system where message delays and differences in relative speed of processors are bounded

The Byzantine Generals Problem 8/14


Algorithm $OM(0)$

1. The general sends his value to every lieutenant
2. Each lieutenant uses the value he receives from the general or uses the default value if he receives no value


Algorithm $OM(t)$, $t > 0$

1. The general sends his value to every lieutenant
2. Let v_i be the value that lieutenant i receives from the general or else be the default value. Lieutenant i acts as the general in algorithm $OM(m-1)$ to send the value to each of the $n-2$ other lieutenants
3. For each i and each $j \neq i$, let v_j be the value that lieutenant i received from lieutenant j in step 2 or else the default value. Lieutenant i uses the value $majority(v_1, \dots, v_n)$


The Byzantine Generals Problem 9/14




The Byzantine Generals Problem 10/14



The Byzantine Generals Problem 11/14



The Byzantine Generals Problem 12/14



The Byzantine Generals Problem 13/14

- OM(m) requires $m+1$ rounds of message exchange
 - No algorithm can reach agreement in less than $m+1$ rounds
- $O(n^m)$ messages
- Use OM(m) to distribute the opinion of each general to other generals in consensus problem

The Byzantine Generals Problem 14/14

- The situation is improved if messages can be *signed*
- General might still be a traitor
- A traitor can't lie about the contents of a message that it received but only fail to pass it on
- Tolerate an *arbitrary number of traitors* but still need $m+1$ rounds of message exchange

References

- [Borg et al. 1989] Anita Borg et al., *Fault Tolerance Under UNIX*, ACM Transactions on Computer Systems, 7(1):1-24, February 1989.
- [Lamport et al. 1982] Leslie Lamport, Robert Shostak, and Marshall Pease, *The Byzantine Generals Problem*, ACM Transactions on Programming Languages and Systems, 4(3):382-401, July 1982.

See also:

- [Jalote 1994] Pankaj Jalote, *Fault Tolerance in Distributed Systems*, Prentice Hall, 1994, ISBN 0-13-301367-7, Chapter 3.