

Atomic multicast 1/4

- Consider implementing atomic multicast

Want:

- Multicast message received by all destinations or none
- Non-delivery only if sender fails during protocol

Assume:

- Implement non-uniform dynamic atomic multicast
 - message received by all destinations that do not fail or none
- Reliable point-to-point communication

Atomic multicast 2/4

Easy!

- Simply have sender send message to all destinations in turn
- What if sender fails while sending?
 - some destinations receive message and others not!

Atomic multicast 3/4

So,

At sender:

- Send message to all sites where there is a destination process

At receiving site:

- If message has not already been received
 - send copy of message to all other sites where there is a destination process
 - deliver message to local destination processes

Atomic multicast 4/4

- This works because if any site receives message and remains operational, all sites will receive the message
 - remember point-to-point links are reliable
- A site will not receive a message only if *no* site that did receive the message stays operational long enough to forward it
- Protocol is expensive because of extra messages sent and need to keep some information around to detect duplicate messages
- Number of messages sent could be reduced if messages are only retransmitted in case of failure of original sender

Ordered multicast

- Will look at totally and causally ordered multicast
- Most ordered broadcast protocols make use of *hold-back*
- A message that is received by some process might be held back (i.e. not delivered) until other messages that should be delivered before it are received and delivered
- Message on the hold back queue may be *deliverable* or *undeliverable*
- A deliverable message may be delivered only when it reaches the head of the queue

Totally ordered multicast 1/9

- Want message delivered to every destination in the same order
- Basic idea is to assign totally ordered timestamps to every message and deliver messages in timestamp order
- Two possible approaches to generating timestamp
 - use a centralised sequencer or token
 - distributed agreement

Totally ordered multicast 2/9

Sequencer

- One process acts as the sequencer
- All messages are sent to the sequencer
- Sequencer assigns a timestamp to each message and multicasts message to all other destinations
- All destinations deliver messages in timestamp order
 - messages held back until all previous messages delivered
- Lost messages are easily detected and can be obtained from sequencer

Totally ordered multicast 3/9

- Simple, but:
- Increased message latency
- Sequencer is single point of failure
- Sequencer may become a bottleneck

Totally ordered multicast 4/9

ISIS total-ordering protocol

- Assume message sent from single sender to multiple destinations
- Each received message is put in the hold back queue of the receiver
 - marked as undeliverable
- The receiver assigns a *proposed* timestamp to the message and returns it in a message to the sender
 - must be larger than any timestamp proposed or received by that process in the past
 - made unique by including process identifier as a suffix to the timestamp
- Sender chooses *largest* proposed timestamp as final timestamp for message and informs destinations

Totally ordered multicast 5/9

- Receivers assign final timestamp to message in hold-back queue and mark message as deliverable
- Hold-back queue is reordered in timestamp order
- When the message at the head of the hold-back queue is deliverable it is delivered

Totally ordered multicast 6/9

Site 1

m3	m1	m2					
15.1	16.1	17.1					
u	u	u					

Site 2

m2	m1	m3					
16.2	17.2	18.2					
u	u	u					

Site 3

m1	m3	m2					
17.3	18.3	19.3					
u	u	u					

Step 1: messages delivered to 3 sites in different orders

m3	m2	m1					
15.1	17.1	17.3					
u	u	d					

m2	m1	m3					
16.2	17.3	18.2					
u	d	u					

m1	m3	m2					
17.3	18.3	19.3					
d	u	u					

Step 2: m1 assigned final timestamp and delivered at site 3

Totally ordered multicast 7/9

m3	m1	m2					
15.1	17.3	19.3					
u	d	d					

m1	m3	m2					
17.3	18.2	19.3					
d	u	d					

m3	m2						
18.3	19.3						
u	d						

Step 3: m2 assigned final timestamp and m1 delivered at node 2

m1	m3	m2					
17.3	18.3	19.3					
d	d	d					

m3	m2						
18.3	19.3						
d	d						

m3	m2						
18.3	19.3						
d	d						

Step 4: m3 assigned final timestamp; all messages delivered

Totally ordered multicast 8/9

Why does this work?

- Guarantees that once a message is delivered no earlier message will arrive
- Notice that because sender always chooses highest proposed timestamp, reordering the hold-back queue can only ever result in a message moving backwards
 - can never be assigned a timestamp lower than its proposed timestamp
- So, when a deliverable message reaches the head of the queue, no message with a lower timestamp can possibly arrive.

Totally ordered multicast 9/9

- Again protocol is expensive:
 - two rounds of communication - similar to 2PC protocol
 - must wait for preceding messages to be delivered
- Can tolerate failure of any destination
- Like sequencer protocol, have to handle failure of sender during protocol execution
- Sender not likely to be a bottleneck

Causally ordered multicast 1/3

- The basic requirement is that before a message is delivered to some process, all the messages that causally precede it must already have been delivered
- Those messages are the ones that were received by the sender of the message before sending the new message
- Could have the sender attach all those messages to each new message
 - by definition, it has them!
- A better solution is simply to attach a description of what those messages were using a vector timestamp

Causally ordered multicast 2/3

- Each process p_i maintains a vector clock VT_i
- VT_i updated as usual as messages are delivered to p_i
 - VT_i encodes the set of messages from all processes that have been delivered to p_i
- VT_i is used to timestamp multicast messages
 - p_i increments $VT_i[i]$ before sending every message and adds it as timestamp TS to the message
- A multicast message can be delivered immediately to its sender
 - by definition, this message follows all the messages that were already delivered to its sender
- At any other process p_j ($j \neq i$) incoming messages are placed on the hold back queue

Causally ordered multicast 3/3

- A message can be delivered only after all messages from the same sender...

$$\text{i.e. } TS[i] = VT_j[i] + 1$$

- ... and all causally preceding messages have been delivered to p_j

$$\text{i.e. } VT_j[K] \geq TS[k] \text{ for all } k \neq i$$

- Moreover, timestamp allows missing messages (and who originally sent them) to be identified
- No single point of failure but need some way of determining when a particular message has been delivered everywhere
- For large groups, size of timestamps is an issue